

It Is Not the Signature: A Pre-Authentication Key Parse Dominates JWT Validation in jsonwebtoken and Is Load-Bearing for Its Key-Confusion Defence

Jannatul Ferdous Katha^{a,c}, Tasmia Tahmid Prova^{b,c}, Md Kishor Morol^c, Tze Hui Liew^{d,*}, Dip Nandi^{c,e}

^a*Department of CSE, Bangladesh University of Professionals (BUP), Dhaka, Bangladesh*

^b*Department of CSE, Bangladesh University of Business and Technology (BUBT), Dhaka, Bangladesh*

^c*ELITE Research Lab, Queens, NY, USA*

^d*Faculty of Information Science and Technology, Multimedia University, Melaka, Malaysia*

^e*Department of Computer Science, American International University-Bangladesh (AIUB), Dhaka, Bangladesh*

Abstract

JSON Web Token (JWT) validation is a common authentication step for microservice requests, and its per-request cost is routinely attributed to cryptography. In `jsonwebtoken`, a widely used Node.js implementation, the signature primitive is 2.208 μ s of a 29.92 μ s validation. The dominant component, 18.04 μ s (60%), is a discarded key-parse attempt: a shared secret supplied as a string is first offered to the asymmetric key parser and is resolved only after that parse fails. The attempt runs before the algorithm check and before signature verification, so any well-formed token, forged or not, triggers it. It entered the library with the security release that fixed an RSA-to-HMAC key-confusion vulnerability, and it is load-bearing: removing it the obvious way fails the library's own key-confusion tests. Its price is set, causally, by the bundled OpenSSL release: the same calls issued from C cost 394.4 μ s on OpenSSL 3.0.16 and 8.447 μ s on 3.5.8. In a containerised service on a single-node Kubernetes testbed, validation adds 384.6 μ s of processor time

*Corresponding author.

Email addresses: `ferdous.katha@elitelab.ai` (Jannatul Ferdous Katha), `tasmia.prova@elitelab.ai` (Tasmia Tahmid Prova), `kishormorol@ieee.org` (Md Kishor Morol), `thliew@mmu.edu.my` (Tze Hui Liew), `dip.nandi@aiub.edu` (Dip Nandi)

per request, within 0.8% of the microbenchmark prediction. Of seven JWT libraries in four languages, only this one attempts the parse. We give a narrow fix that preserves the key-type defence and makes validation 4.15 times faster, show that a client controlling the token header can still reach the parse under it, and identify pre-parsed keys as the remedy that removes it for every input. In our sample, 94% of call sites pin no algorithm and none pre-parses the key. The defect has been reported upstream.

Keywords: JSON Web Token, authentication, key confusion, cryptographic API, OpenSSL, Node.js, performance measurement

1. Introduction

Authentication in cloud-native services is commonly treated as a per-request cryptographic cost, and measurement of layered enforcement has concentrated on the transport and proxy layers where that cryptography runs (Zhu et al., 2023). The assumption is old and has been tested before, one layer down: Coarfa et al. (2006) instrumented TLS web servers and found RSA responsible for only a minority of server cost, the remainder falling to symmetric operations, hashing and protocol handling. Our result is that argument repeated at the application layer, and reaching a sharper version of the same conclusion.

The reasoning being tested is intuitive: validating a JSON Web Token (Jones et al., 2015) means verifying a signature, signatures are cryptography, and cryptography is expensive relative to parsing. The conclusion usually drawn is that the cost should be *moved* (to a sidecar proxy, a gateway, or dedicated hardware) because specialised components do cryptographic work well.

We report that the premise is wrong for `jsonwebtoken`, a widely used JavaScript implementation. A second natural explanation fails as well. Once the signature primitive is shown to be a small share of the cost, the remainder is readily attributed to key material being converted from a string into a key object on every call. Measuring that conversion directly puts it at 0.583 μ s: too small to be the explanation for anything.

The actual mechanism is a control-flow accident, shown in Figure 2. Given a secret that is not already a key object, the library attempts `createPublicKey()` first and falls back to `createSecretKey()` in the exception handler. For an HMAC secret the first call cannot succeed. Every

validation therefore constructs and discards an OpenSSL parse failure, at 18.04 μ s against a total call cost of 29.92 μ s.

That 18.04 μ s is the cost of the whole discarded attempt: the work OpenSSL performs before deciding the input is not an asymmetric key, plus the construction and unwinding of the exception that carries that decision back to the library. Section 4.2 separates the two, by reissuing the OpenSSL calls from C and by measuring an exception of the same shape with no OpenSSL call. Where the attempt is expensive it is almost all parse; the exception is a few microseconds.

How much the attempt costs depends on which OpenSSL release performs it, by more than an order of magnitude (Section 5). Runtimes carrying an OpenSSL release from 3.2 onward, which include the current Node.js release lines, pay tens of microseconds per validation; runtimes carrying the 3.0 series pay hundreds. The 3.0 series is not confined to retired Node.js releases. It is the system OpenSSL of long-term-support distributions still in wide use, among them Ubuntu 22.04 and Debian 12, and a Node.js build that links the system library rather than its bundled copy inherits it; so does any service whose image is pinned to an older base layer, which is common in practice (Zerouali et al., 2019). At either price the discarded attempt is the largest single component of the call, and in every case it is avoidable.

The cost is also a security matter, in three ways (Section 8). First, it is paid before authentication: key resolution precedes both the algorithm check and the signature check, so every request carrying a well-formed token pays it, whether or not anyone the service trusts issued that token. Second, it did not arrive through neglect. The resolution block entered the library in version 9.0.0, the release that fixed an RSA-to-HMAC key-confusion vulnerability (Auth0, 2022), and the key-type checks that implement that fix read the type of the key object the block produces. Third, it therefore cannot be removed the obvious way: dispatching on the token’s declared algorithm without restricting which key material may take the fast path fails the library’s own key-confusion tests. A performance change to authentication code is a change to its security, and we analyse it as one.

The result is broader than one library in two respects. The design that produces the cost (resolving key material on every call, by way of a parse whose failure serves as a type test) can be named independently of `jsonwebtoken`, and Section 7 shows it predicts which of seven libraries pay. And the reason the cost went unattributed is methodological: the component that sets its price, the bundled OpenSSL release, is one that measurement protocols rarely

65 record at the granularity at which it changes.

66 *Research questions.*

67 **RQ1** Where does the per-request cost of token validation actually go?

68 **RQ2** What governs the magnitude of that cost?

69 **RQ3** Does it matter in a running containerised service, and how common is
70 the pattern that avoids it?

71 **RQ4** Is the defect a property of one library, or of a design choice other
72 libraries could share?

73 **RQ5** How does the cost relate to the library’s security: who can cause it, why
74 is it there, and can it be removed without weakening the key-confusion
75 defence?

76 *Contributions.* All supported by measurements in the replication pack-
77 age (Katha et al., 2026).

- 78 1. A decomposition of the validation call, by two independent methods,
79 locating its dominant component in a discarded key-parse attempt
80 rather than in the signature primitive or in key conversion, and a
81 second decomposition of that attempt into the OpenSSL parse, the V8
82 exception and the binding between them (Section 4).
- 83 2. A security analysis of that component (Section 8). It is reachable before
84 authentication by any well-formed token; it entered the library with a
85 key-confusion fix and is load-bearing for it, so the natural optimisation
86 fails the library’s key-confusion tests; a narrow fix preserves the defence
87 and makes validation 4.15 times faster, but leaves the parse reachable
88 through an attacker-chosen algorithm header; and only a pre-parsed
89 key removes it for every input.
- 90 3. A causal attribution of its magnitude to the bundled OpenSSL release:
91 the parse reissued from C with no JavaScript engine across 8 identically
92 built OpenSSL releases, a runtime pair that crosses engine and library,
93 and a repetition on a second architecture (Section 5).
- 94 4. A measurement in a containerised service across seven arrival rates, and
95 a survey establishing that the avoiding pattern is rare in three samples
96 of public code (Section 6).
- 97 5. A comparison of seven JWT libraries in four languages showing that
98 the cost follows a design choice (resolving key material on every call,
99 of which a failing asymmetric parse used as a type test is the costliest
100 form) rather than JWT validation as such (Section 7).

101 2. Background and related work

102 *What a validation call does.* A JSON Web Token carries a header, a payload
103 and a signature over the first two (Jones et al., 2015). Validating one means
104 decoding two segments, resolving key material into a form the cryptographic
105 library accepts, and verifying the signature. The middle step, key resolution,
106 is the subject of this paper, and it is the step most often left implicit when
107 validation cost is discussed.

108 *Where the cost is usually attributed.* Published measurement of layered en-
109 forcement in microservices concentrates on transport and proxy layers. Zhu
110 et al. (2023) decompose service-mesh sidecar overhead and show that the dom-
111 inant contributors depend on configuration and workload rather than being a
112 fixed tax; we reach an analogous conclusion one layer up, for application-level
113 validation. Application-level token validation is itself thinly evidenced, and is
114 commonly reported as a difference between scenario averages, which attributes
115 to validation every difference between two runs. We were unable to locate
116 a study that measures the key-resolution step directly; we state that as an
117 absence we searched for rather than as a claim about the whole literature.

118 *Cryptographic cost has been over-attributed before.* Coarfa et al. (2006) remains
119 the clearest precedent: measuring TLS web servers rather than primitives,
120 they found the public-key operation was a minority of server cost. Work on
121 cryptographic library performance since has largely stayed at the level of
122 the primitive or the protocol exchange (Naser et al., 2020; Sikeridis et al.,
123 2020), with Tuveri and Brumley (2019) a notable exception in examining how
124 a library loads and dispatches implementations rather than how fast those
125 implementations run. Our result concerns neither: it is the cost of getting
126 key material *into* the library.

127 *Performance bugs as a class.* Jin et al. (2012) studied real-world performance
128 bugs across five code bases and found that most follow a small number of
129 recurring patterns (among them uncoordinated and skippable function calls)
130 and that fixes are usually local. Zaman et al. (2012) found, from 400 Firefox
131 and Chrome reports, that performance bugs differ from other bugs in how
132 they are reported, discussed and fixed, in part because reproducing them is
133 hard; Nistor et al. (2013a) report that most are found by code reasoning rather
134 than by observing symptoms. Liu et al. (2014) reach similar conclusions for
135 mobile applications. Detection work targets structural symptoms: redundant

136 traversals with similar memory-access patterns (Nistor et al., 2013b), or
137 predicates that separate slow from fast runs (Song and Lu, 2014). The defect
138 here is of the “skippable call” kind Jin et al. describe, but it has none of the
139 structural signatures those detectors look for: it is one call, executed once
140 per request, whose failure is its only observable effect. It is also the shape
141 that makes a performance bug hard to report: our own upstream report was
142 preceded by a two-year-old report of the symptom with no diagnosed cause
143 (Section 8.6).

144 *JavaScript and Node.js performance.* Selakovic and Pradel (2016) classified
145 performance issues in JavaScript projects and found that inefficient API usage
146 (correct calls made in a costly way) is among the most common root causes,
147 and that the fixes are typically a few lines. Their later profiler (Selakovic et al.,
148 2017) targets precisely the *order* in which a program evaluates conditions, the
149 axis on which this defect lives: a check that almost always fails is evaluated
150 before one that almost always succeeds. Gong et al. (2015) locate code that
151 defeats JIT optimisation; exception-heavy control flow is one such pattern,
152 and Section 4.2 measures what it costs here. Work on the Node.js event loop
153 shows why a per-request synchronous cost matters beyond its magnitude: a
154 single slow handler delays every other client (Davis et al., 2018; Staicu and
155 Pradel, 2018). We make no availability or capacity claim, but the mechanism
156 by which a synchronous per-request cost propagates is the same.

157 *Version and environment variation.* That software performance changes
158 across versions in ways maintainers do not anticipate is established: Kalte-
159 necker et al. (2023) track performance across configurable systems’ release
160 histories, and Chen and Shang (2017) study the code changes that introduce
161 regressions. Laaber and Leitner (2018) assess whether microbenchmark suites
162 detect such changes at all, and Stefan et al. (2017) find performance unit
163 testing rare in practice. Container images make this concrete by pinning a
164 dependency set implicitly: Cito et al. (2017) characterise how the ecosystem
165 uses them, Zerouali et al. (2019) show that images routinely ship outdated
166 packages, and Malka et al. (2026) show that an image tag does not guarantee
167 a reproducible build. For OpenSSL 3 specifically, the maintainers of a major
168 downstream library have stated publicly that parsing and key loading re-
169 gressed substantially against 1.1.1 (Gaynor and Kehrer, 2026); that statement
170 is not peer reviewed and was not produced under a protocol we can assess,
171 and Section 5.5 supplies a controlled measurement for the one path this paper
172 concerns. Section 5 is an instance with a measured cost attached.

173 *Measurement protocol is the thing at stake.* The methodological core of this
174 paper is not new. Georges et al. (2007) established that a single execution
175 of a managed-runtime benchmark is not a measurement, and that intervals
176 must be computed across executions rather than within one; Kalibera and
177 Jones (2013) show how to spend a fixed budget across levels of repetition,
178 which is the reasoning behind our choice of many short processes over few
179 long ones. Mytkowicz et al. (2009) showed that environmental details unre-
180 related to the code under study can produce effects larger than the one being
181 measured, without any obvious mistake, and Curtsinger and Berger (2013)
182 that memory layout alone can do so. Barrett et al. (2017) found that virtual
183 machines frequently do not reach a steady state, Traini et al. (2023) that
184 developers misjudge when warmup ends, and Costa et al. (2021) catalogue how
185 microbenchmarks go wrong in practice. Mytkowicz et al. (2010) showed that
186 sampling profilers disagree on which method is hot, which is why Section 4
187 uses the profiler only as corroboration of a result obtained by differencing and
188 relies on ordering rather than on its percentages. Our contribution is not to
189 these methods but is an instance of what they warn about. The attribution
190 this paper overturns has exactly the structure Mytkowicz et al. describe: a
191 second variable, the bundled OpenSSL version, moving unrecorded alongside
192 the one under study, at a granularity coarse enough to hide it (Section 5).

193 *Environment as a confounder.* Laaber et al. (2019) quantify how far cloud
194 environments degrade microbenchmark reliability, and Leitner and Cito (2016)
195 characterise performance variation across public infrastructure providers.
196 Their concern is variance between nominally identical environments. Ours
197 is adjacent and, we think, less examined: a *systematic* difference between
198 environments that are not identical and are routinely treated as though they
199 were: two container base images differing in a bundled library version. Laaber
200 et al. are also the reason the second-architecture repetition in Section 5.6 is
201 reported as an ordering check rather than as a replication of magnitudes.

202 *Cryptographic interfaces are hard to use correctly.* The defect reported here
203 is not a cryptographic weakness but a consequence of how a cryptographic
204 interface is invoked, which is a well-documented class. Lazar et al. (2014)
205 found that most of the cryptographic vulnerabilities they studied were misuses
206 of libraries by applications rather than bugs in the libraries. Egele et al. (2013)
207 found cryptographic misuse widespread in deployed applications and Georgiev
208 et al. (2012) found certificate validation broken across a wide range of non-
209 browser software; Nadi et al. (2016) attribute much of it to the interfaces

210 themselves rather than to developer inattention, a position Green and Smith
211 (2016) put directly and Indela et al. (2016) respond to with interface proposals;
212 and Acar et al. (2017) show experimentally that API design materially changes
213 whether developers produce correct code. Detection has been automated by
214 specifying correct usage, as CrySL does (Krüger et al., 2018), and by static
215 analysis at scale (Rahaman et al., 2019). Our finding sits within that line
216 with one difference: the interface here admits a use that is *correct*, and costly
217 only for a reason invisible at the call site. A CrySL-style rule could express
218 “pass a key object, not a string”, but the rule sets we know of encode security
219 properties only, so none would report it.

220 *JWT and JOSE libraries.* Security analyses of JWT implementations consis-
221 tently find key-handling defects. Detering et al. (2017) demonstrated attacks,
222 key confusion among them, against several popular JOSE libraries; Xu et al.
223 (2023) automated the detection of cryptographic vulnerabilities in applica-
224 tions that use JWT; Yang et al. (2026) evaluated 43 JWT libraries across
225 16 programming languages for vulnerabilities; and Shatnawi et al. (2024)
226 surveyed Python JWT libraries for security warnings and adoption. RFC
227 8725 (Sheffer et al., 2020) codifies the resulting practice. All of this work
228 treats a library’s handling of key material as a security surface. Section 7
229 treats the same surface as a cost site, across seven libraries, and Section 8.4
230 shows the two views are coupled: the control flow that costs time in one
231 library is the control flow that enforces its key-type check.

232 *Mining public repositories.* Section 6.5 draws on code hosted publicly, and
233 the hazards of doing so are catalogued by Kalliamvakou et al. (2014), among
234 them that many repositories are not software projects, that activity is not
235 uniformly distributed, and that a search result is not a random sample. Our
236 sampling frame is stated in Section 3.10 with those hazards in view, which is
237 why no population proportion is computed from it.

238 *The cost of exceptions.* That exceptions are expensive in managed runtimes,
239 and that their cost is concentrated in construction rather than in unwinding,
240 is long established: Ogasawara et al. (2006) optimise around exception paths
241 in a Java virtual machine, and Renwick et al. (2019) quantify the cost in a
242 compiled setting. Parravicini and Mueller (2021) provide the closest available
243 cost-attribution baseline for the V8 engine, though on speculation rather than
244 exceptions. Empirical studies of how exceptions are *used* find handlers that
245 discard what they catch to be common: Cabral and Marques (2007) report

246 that handlers in the Java and .NET code they studied rarely attempt recovery,
247 Kery et al. (2016) find that handlers mostly log, print, return or rethrow, and
248 de Pádua and Shang (2017) measure the prevalence of handling anti-patterns,
249 as Weimer and Nacula (2004) did earlier for error-handling mistakes. That
250 literature treats a swallowed exception as a correctness and maintainability
251 smell. The handler at the centre of this paper is a swallowed exception used
252 as control flow on a hot path, and it is a performance defect rather than
253 a correctness one: the fallback it guards is correct. None of these studies
254 measures such a handler’s cost; Section 4.2 does, and finds the exception itself
255 the smaller part.

256 *Key-type confusion.* That a token declaring an HMAC algorithm must not be
257 validated against asymmetric key material is long established; RFC 8725 (Shef-
258 fer et al., 2020) describes the attack and its mitigations. The library studied
259 here has been hardened against this class twice. Versions before 4.2.2 accepted
260 an HMAC token verified with an RSA or EC public key used as the shared
261 secret (CVE-2015-9235) (GitHub Advisory Database, 2015), and versions
262 up to 8.5.1 could still be led to verify forged tokens from RSA to HMAC
263 through their key-retrieval path (CVE-2022-23541), alongside two related
264 defects in key-type and default-algorithm handling (CVE-2022-23539, CVE-
265 2022-23540) (Auth0, 2022); all three were fixed in version 9.0.0. Section 8
266 reports no new attack. It reports that the second hardening introduced the
267 control flow whose cost this paper measures, that the defence depends on
268 that control flow, and that removing the cost naively removes the defence.

269 2.1. What we searched for and did not find

270 Table 1 summarises the positioning. Each cell records only what this
271 section states about the work cited; it is a map of the lines of work this paper
272 connects, not an assessment of them.

273 Three areas bear directly on this result and we were unable to locate
274 verifiable peer-reviewed work in any of them. We state each as an absence
275 we searched for rather than as a claim about the whole literature; each is
276 also, in part, why the misattribution corrected here survived as long as it
277 did. Two adjacent questions (the cost of error construction in V8, and where
278 key-parsing cost lies inside OpenSSL) were equally unstudied, and this paper
279 measures both directly (Sections 4.2 and 5.5).

280 *OpenSSL 3.x performance regressions.* We found no peer-reviewed study
281 of performance regressions across the OpenSSL 3.x series. Evidence exists

Table 1: Positioning against the closest lines of work. *JWT*: studies JWT or JOSE implementations. *Crypto-API security*: treats how a cryptographic interface or key material is used as a security concern. *Cost*: studies a performance cost or performance defect. *Version as cause*: attributes a cost or defect to the version of a bundled dependency. *Cost-defence coupling*: shows that removing a cost would weaken a security defence. ✓: addressed; ◦: partly (see note); -: not addressed.

Work	JWT	Crypto-API security	Cost	Version as cause	Cost-defence coupling
Detering et al. (2017); Xu et al. (2023); Yang et al. (2026)	✓	✓	–	–	–
Egele et al. (2013); Nadi et al. (2016); Krüger et al. (2018); Rahaman et al. (2019)	–	✓	–	–	–
Coarfa et al. (2006); Naser et al. (2020); Sikeridis et al. (2020)	–	–	✓	–	–
Zhu et al. (2023)	–	–	✓	–	–
Jin et al. (2012); Selakovic and Pradel (2016)	–	–	✓	–	–
Kaltenecker et al. (2023); Chen and Shang (2017)	–	–	✓	◦ ^a	–
Zerouali et al. (2019); Cito et al. (2017)	–	–	–	◦ ^b	–
This paper	✓	✓	✓	✓	✓

^aVersions of the system under study, not of a bundled dependency. ^bContainer images pin outdated dependency versions; no cost is measured.

as a maintainer statement (Gaynor and Kehrer, 2026), as project issues (for instance `openssl/openssl` issues #17627 and #20698, which concern performance of the 3.x provider and decoder machinery) and as mailing-list discussion. None was produced under a measurement protocol we can assess, so Section 5.5 measures the path directly rather than adopting the mechanism those sources propose.

Base image to bundled library to application performance. We found no academic study tracing that chain. Malka et al. (2026) address build reproducibility, Zerouali et al. (2019) package staleness and Cito et al. (2017) ecosystem practice; none measures how the library version an image pins changes application performance.

Cryptographic API misuse as a performance defect. The misuse literature




Figure 1: Study workflow. Top: the five research questions. Middle: the instruments that answer each, with the section describing it. Bottom: the tables and figures each produces. The C probe and the exception baseline (RQ2) also decompose the discarded attempt for RQ1 (Table 3), and the call-site survey (RQ3) supplies the algorithm-pinning figure used in RQ5. All measurements feed one replication package, from which every reported value is generated.

294 cited above treats misuse exclusively as a security problem. We found no
 295 work treating a correct-but-costly use of a cryptographic interface as a latency
 296 or throughput defect, which is precisely what this paper reports.

297 3. Methodology

298 Figure 1 gives the study’s workflow: the instruments used for each research
 299 question and the tables and figures each produces. Every instrument writes
 300 to committed run files, from which every number in the paper is generated
 301 (Data availability).

302 3.1. Testbed

303 Unless stated otherwise, measurements were taken on one Apple-silicon
 304 machine running macOS 26.5.2 (arm64), referred to throughout as *the*
 305 *host*. The host microbenchmarks of Section 4 ran directly on macOS un-
 306 der Node.js v26.6.0. Container measurements (Sections 5 and 6) ran in
 307 the Linux/arm64 virtual machine of Docker Desktop’s built-in single-node
 308 Kubernetes cluster (v1.32.2).

309 The service studied in Section 6 is a Node.js microservice (Express 5.1.0,
 310 `jsonwebtoken` 9.0.3) built on the `node:18-alpine` image, run as a single
 311 replica with an Istio sidecar (`istio/proxyv2`; the proxy version was not
 312 recorded, and its image digest is in the replication package), and reached
 313 over TLS through an NGINX ingress controller that is itself injected into
 314 the mesh. It verifies HS256 tokens with a shared secret supplied as a string.
 315 Load was generated in open loop by k6 v2.2.0 on the same machine, so
 316 the load generator, the ingress, the sidecars and the service share one host
 317 (Section 3.11). This is a controlled single-node testbed, not a production
 318 deployment; we call it the *containerised service*. The x86_64 measurements
 319 used the separate machine described in Section 3.8.

320 3.2. Unit of measurement

321 The unit is one operating-system process, not one call. Each process
 322 performs 40000 timed iterations of a single condition after a warmup of half
 323 that number, and contributes its median. Reported figures are medians of
 324 per-process medians, with percentile bootstrap intervals taken over processes.

325 An interval computed over the calls inside one process describes that
 326 process’s steady state and is far too narrow: it cannot observe the compilation
 327 decisions, inline-cache states and heap layouts that differ between runs of the
 328 same code (Georges et al., 2007). Conditions in the decomposition of Section 4
 329 were run in 15 separate processes of 40000 calls; conditions in the runtime
 330 matrix of Section 5 in 10 processes of 20000 calls. The largest between-process
 331 coefficient of variation is 3.5% on the host; in the matrix it is 20.2%, occurring
 332 on the `probe_succeeds` condition under `node:26.6.0-bookworm` rather than
 333 across the matrix generally.

334 Bootstrap intervals over 15 or 10 units are coarse, and we report them as
 335 indicative rather than as calibrated coverage. Where an interval collapses to
 336 a point (the timer overhead row of Table 2), it should be read as the clock’s
 337 resolution showing through, not as a precise estimate. The per-process values
 338 underlying every interval are in the replication package.

339 Conditions are interleaved within each round rather than run in blocks.
 340 Under block ordering a thermal excursion or a background process lands
 341 entirely on whichever condition is executing and is indistinguishable from an
 342 effect; interleaved, it is distributed across conditions and appears as between-
 343 round variance. The clock’s own cost is measured in the same process as the
 344 condition it accompanies and is reported rather than subtracted: 0.042 μ s
 345 against the smallest quantity reported, 0.583 μ s.

346 Differences between conditions are reported throughout as the median
347 of per-round *paired* differences, not as the difference of two medians. The
348 two are not identical and the distinction matters at the magnitudes involved;
349 Section 4.3 and Section 8.4 each rely on it.

350 3.3. Decomposing the call

351 `jwt.verify()` is decomposed into the operations it performs, each timed
352 in isolation under the same protocol: structural decode; the HMAC primitive
353 with a string key and with a pre-parsed key; the symmetric key constructor;
354 and the asymmetric key parse the library attempts first, in both the case
355 where it throws and the case where it succeeds.

356 The library measured throughout is `jsonwebtoken` version 9.0.3, cloned
357 at tag v9.0.3. The code path at issue is shown in Figure 2.

```
120 if (secretOrPublicKey != null && !(secretOrPublicKey instanceof
    KeyObject)) {
121   try {
122     secretOrPublicKey = createPublicKey(secretOrPublicKey);
123   } catch (_) {
124     try {
125       secretOrPublicKey = createSecretKey(typeof
secretOrPublicKey === 'string' ? Buffer.from(
secretOrPublicKey) : secretOrPublicKey);
126     } catch (_) {
127       return done(new JsonWebTokenError('secretOrPublicKey is not
valid key material'))
128     }
129   }
130 }
```

Figure 2: Key resolution in `jsonwebtoken` version 9.0.3, `verify.js`, lines 120--130. Verbatim but for the common leading indentation, which is removed to fit the column; line numbers are the file's own. A secret that is not already a `KeyObject` is offered to `createPublicKey()` first, and `createSecretKey()` is reached only from the exception handler. For an HMAC secret supplied as a string the outer `try` can never succeed, so the asymmetric parse is attempted and discarded on every validation.

358 3.4. *Separating the runtime from the deployment*

359 A host microbenchmark compared against an in-container measurement
360 cannot support an attribution to deployment: those two environments differ
361 in more than deployment, and any of the differences could carry the effect.
362 Here containerisation is held constant and the runtime is varied instead: four
363 container images spanning two OpenSSL series and three V8 major versions,
364 on one machine under one container engine.

365 The pair carrying most of the inference is Node 18 and Node 20: they ship
366 V8 10 and V8 11 respectively while sharing the OpenSSL 3.0 series. In stock
367 images Node and OpenSSL versions otherwise co-vary, which would leave the
368 attribution underdetermined. Because that pairing carries the attribution,
369 the runtime identity of each image (Node, OpenSSL and V8 version, libc,
370 architecture and image digest) is recorded as a datum and cross-checked
371 against the versions each measurement row carried. That check found 0
372 mismatches across the matrix. Section 5.2 states what this design alone
373 cannot separate, and Sections 3.5 and 3.6 describe the two further designs
374 that do.

375 3.5. *Taking the JavaScript engine out*

376 Runtime pairs can separate engine from library only as far as stock
377 builds allow, so the attribution was tested by removing the engine al-
378 together. We read the Node source at the two release lines measured
379 (`src/crypto/crypto_keys.cc` at v18.20.8; `deps/ncrypto` at v22.23.3,
380 which issues the same sequence) and reproduced in C, call for call, what
381 `createPublicKey()` asks of OpenSSL when handed an HMAC secret: three
382 `PEM_bytes_read_bio` scans for a public key, an RSA public key and a
383 certificate, each under an error-queue mark and each followed by `BIO_reset`;
384 then, since none was recognised, `PEM_read_bio_PrivateKey` with Node's
385 passphrase callback; then the C side of Node's error path, which drains
386 the OpenSSL error queue into strings. The V8 side of that error path is
387 deliberately absent; it is what Section 3.6 measures.

388 The program was compiled once per OpenSSL release with one compiler
389 (GCC 13.3.0) and one set of flags, against 8 releases from 3.0.16 to 3.6.1, each
390 built from source with identical configuration and linked by run path so that
391 no system library could be picked up. Each binary reports the version it
392 actually loaded; rows whose loaded version differed from the one requested
393 would be discarded, and there were 0. The protocol is that of Section 3.2:
394 one condition per process, 15 processes per release and condition, 5000 timed

395 calls each, timed with `clock_gettime`, conditions and releases interleaved
396 within rounds. Besides the full sequence we time its parts separately (the
397 three public-key scans, the private-key parse, the error-string drain) and a
398 *successful* parse of an RSA-2048 public key through the same path.

399 3.6. *Separating the exception from the parse*

400 Inside Node the failed parse is an OpenSSL failure followed by a V8
401 exception: Node constructs an `Error`, attaches `library`, `reason` and `code`
402 properties, throws it through the `createPublicKey()` frame, and the caller
403 discards it. To measure that exception without the parse, a JavaScript
404 benchmark builds the same object (message, property values and error code
405 are captured from a real failure on the runtime being measured, not typed)
406 and throws it from 8 frames below the catch. Variants separate its parts: a
407 plain `Error` thrown and caught in one frame; the Node-shaped error at depth;
408 the same with stack capture disabled; and a Node-internal argument error
409 raised before any OpenSSL call. The real probe is timed alongside, with and
410 without stack capture, so that the stack’s share of it is obtained by paired
411 differencing. All conditions run under official Node release binaries, which
412 bundle their own OpenSSL, on six runtimes (18.20.8, 20.20.2, 22.23.3, 23.11.1,
413 24.21.0 and 26.10.0) at 15 processes of 10000 calls per condition.

414 Combining this with Section 3.5 gives a three-way decomposition of the
415 probe on each runtime: the C cost of the OpenSSL release that runtime
416 bundles, the V8 exception, and a remainder we attribute to the binding
417 between them: argument preparation, key-handle allocation and the C++
418 side of error construction. The three come from different processes and cannot
419 be paired, so the interval on the remainder is obtained by resampling each
420 run’s processes independently.

421 3.7. *Other libraries*

422 To ask whether the defect is one library’s or a design choice’s, we read the
423 key-handling code of six further libraries, chosen as the most used in each
424 ecosystem: `fast-jwt` and `jose` (JavaScript), `PyJWT` (Python), `golang-jwt`
425 (Go), and `nimbus-jose-jwt` and `jjwt` (Java). For each we answer one
426 question from source: given a string HMAC secret, does it attempt to parse
427 it as an asymmetric key? We then time HS256 verification in each with the
428 secret in two forms, under the same one-condition-per-process protocol at 15
429 processes of 20000 calls. “String” is the form each library’s documentation
430 shows for an HMAC secret, used the way its interface is meant to be used;

431 “pre-parsed” is the library’s own key object, built once. Where a library
432 does not accept a string we use its rawest accepted form and say so. The
433 JavaScript libraries are run under two runtimes, one bundling OpenSSL 3.0
434 and one bundling 3.5, because the `jsonwebtoken` penalty depends on that.
435 The penalty is the median of per-round paired differences between the two
436 forms.

437 3.8. A second architecture

438 Tables 2 and 4 were measured on one ARM64 host. Both were repeated
439 with unchanged scripts on an x86_64 machine (Intel(R) Xeon(R) Processor @
440 2.80GHz, four virtual CPUs), including the four container images of Table 4;
441 Sections 3.5–3.7 were measured on the same machine. That machine is a
442 virtual machine on shared cloud infrastructure, not the dedicated host a
443 replication of magnitudes would require (Laaber et al., 2019; Leitner and Cito,
444 2016). We use it for what it can support: whether the *ordering* of conditions
445 and runtimes holds on a second instruction set, and ratios between conditions
446 measured interleaved in the same rounds. We do not compare its absolute
447 magnitudes with the ARM64 host’s.

448 3.9. In-situ measurement

449 The containerised service (Section 3.1) exposes a histogram timing the
450 verification call alone. The instrument is a *paired-scrape difference*: the
451 cumulative histogram is read before and after each arrival-rate step, and the
452 two readings are differenced, which isolates that step’s distribution from the
453 cumulative counters. Means are computed as the difference of the histogram’s
454 `_sum` divided by the difference of its `_count`, so they are exact and do not de-
455 pend on bucket boundaries; bucket resolution bounds the reported percentiles
456 only.

457 Seven arrival rates are driven, and the whole sweep is run twice on one
458 process: once in ascending order and once in descending order. We report
459 the *range across the two passes*. This is a range of two, not a confidence
460 interval; no standard error is computed from it, and it is not written as a
461 symmetric tolerance. The opposite-order design also serves as a drift test,
462 which Section 6.2 reports.

463 One process and two passes is the whole of the in-situ evidence. It is
464 sufficient to establish that the cost is present in the containerised service and of
465 what order; it is not sufficient to characterise the shape of the cost-versus-rate
466 relation, and Section 6.1 is written accordingly.

467 3.10. *Prevalence: three instruments, three frames*

468 Prevalence is approached with three instruments that fail in different ways.
469 Their sampling frames differ and are not interchangeable, so no population
470 proportion is computed from them. All three were captured on 2026-09-18.

471 *Population co-occurrence counts.* Code-search counts over public repositories, in four cells: two languages by two import forms. The denominator of a
472 cell is files importing the library; the numerator is files that *also mention* the
473 pre-parsing interface. This measures co-occurrence of two strings in one file,
474 not verified usage. The denominators are returned rounded (each is an exact
475 multiple of 1024), so they are magnitudes rather than counts established file
476 by file.

478 *A classified sample of call sites.* Call sites located by search, deduplicated
479 to one file per repository so that no project dominates, with the key argument
480 traced within its file. This is a convenience sample from a relevance-ranked,
481 undocumented ranking; inclusion probabilities are unknown and no weighting
482 back to the population is possible.

483 *A targeted counter-search.* Every file co-occurring the library with the
484 pre-parsing interface, fetched in full and adjudicated by hand. This frame
485 is deliberately biased toward finding the pattern (that is its purpose as a
486 negative control) and it is therefore disqualified as a prevalence estimator.

487 Three reasons preclude a population-level prevalence. The population
488 numerator and the verified numerator measure different things: of 13 auto-
489 matically flagged call sites, 5 were false positives on hand inspection, so the
490 co-occurrence signal carries an error rate that would be inherited, unmeasured,
491 by any extrapolation. The sample has no inclusion probabilities. And the
492 counter-search is selected on the outcome. Results are therefore reported
493 *in our sample*, or *among repositories that co-occur the interface and call the*
494 *validator*, and never as a population proportion. This constraint is carried
495 into the abstract and the conclusion.

496 3.11. *Host state is recorded per measurement*

497 A service measured on a host it shares with its own load generator,
498 ingress and sidecars is measured together with whatever else that host is
499 doing. We observed this directly while building the instrument: re-running
500 one configuration with a background indexing process consuming host CPU
501 moved the point at which throughput collapsed by hundreds of requests per
502 second, and both runs are committed in the replication package (Katha et al.,
503 2026). A curve that moves with unrelated host activity does not characterise

the service, and this paper therefore makes no capacity claim. Every in-situ window instead carries the host load average observed during it, committed beside the measurement, and Section 6.3 uses it.

4. RQ1: Where the cost goes

4.1. The decomposition

Table 2 gives each operation in the validation path, timed in isolation on the host under the protocol of Section 3, and Figure 3 plots it.

Table 2: Cost of each operation in the validation path. Medians of per-process medians over 15 processes of 40000 calls; 95% percentile bootstrap intervals over processes, which with 15 units are indicative rather than calibrated (Section 3.2).

Operation	Median (μs)	95% interval (μs)
<code>verify()</code> , string secret	29.92	[29.75, 30.12]
discarded asym. parse	18.04	[17.83, 18.25]
successful asym. parse	13.83	[13.62, 13.92]
<code>verify()</code> , pre-parsed key	6.042	[5.958, 6.125]
HMAC, string key	2.208	[2.208, 2.250]
HMAC, pre-parsed key	2.166	[2.125, 2.167]
structural decode	1.292	[1.250, 1.292]
symmetric key constructor	0.583	[0.542, 0.583]
timer overhead	0.042	[0.042, 0.042]

The two emphasised rows are the result. The signature primitive costs 2.208 μs of a 29.92 μs call. The conversion to which the remainder is most readily attributed costs 0.583 μs . The discarded attempt that precedes that conversion costs 18.04 μs , which is 60% of the whole call and its largest single component.

Supplying a pre-parsed key instead of a string reduces the call by 23.92 μs [23.58, 24.00] over 15 paired rounds.

Key conversion does not explain it. Once the signature primitive is shown to be small, one explanation for the remainder suggests itself immediately: that the cost lies in converting key material from a string into a key object on every call. It is the natural explanation, and Table 2 refutes it. That




Figure 3: (a) Every operation in Table 2, on a linear scale, with the two components the paper turns on emphasised; whiskers are the 95% percentile bootstrap intervals over 15 processes. (b) The 23.92 μs penalty for supplying a string secret rather than a pre-parsed key, decomposed. The three segments are medians of per-round paired differences, not a partition of a single measurement, and they do not sum to the penalty: the 5.292 μs remainder is unattributed (Section 4.3).

conversion costs 0.583 μs , so the attribution is wrong by roughly the ratio between that figure and 18.04 μs .

A second observation is usually offered in support of the conversion account: an asymmetry between HMAC and RSA costs, read as conversion being dearer than parsing a PEM. It is backwards. A successful asymmetric parse costs 13.83 μs against 0.583 μs for the conversion. The real asymmetry is that the RSA path’s parse *succeeds*, making its cost useful work, while the HMAC path’s fails, making it waste.

4.2. What the 18.04 μs is made of

The 18.04 μs is the cost of the discarded attempt taken whole. It contains at least two things: the work OpenSSL performs before concluding that the input is not an asymmetric key, and the construction and unwinding of the exception that returns that conclusion to the library. Table 2 already bounds the split: a *successful* asymmetric parse (the same entry point, reaching a conclusion and throwing nothing) costs 13.83 μs , the same order as the failing one, so the exception cannot be most of it. Table 3 replaces the bound with a decomposition, combining the C measurement of the same OpenSSL release (Section 3.5) with the exception baseline (Section 3.6) on six official runtimes.

Three things follow. First, on every runtime bundling OpenSSL 3.0 the C parse alone matches the probe to within a few percent: 101% of it under

Table 3: The discarded attempt decomposed, per runtime, in μs (x86_64). *Probe*: `createPublicKey(secret)` inside Node, thrown and caught. *C parse*: the same OpenSSL calls issued from C against the OpenSSL release that runtime bundles, built from source. *V8 exception*: an `Error` of the same shape, message and properties thrown from 8 frames and discarded, with no OpenSSL call. *Binding*: the probe less the other two; its interval, from independent resampling of the three runs, is in the replication package. Medians of per-process medians over 15 processes (probe, exception) and 15 processes (C). The C column uses a stock build of each release, not Node’s own build of it, so the binding column is approximate (Section 10).

Runtime	OpenSSL	Probe	C parse	V8 exception	Binding
Node 18.20.8	3.0.16	390.5	394.4	4.781	-8.694
Node 20.20.2	3.0.19	443.4	431.3	4.796	7.292
Node 23.11.1	3.0.16	383.5	394.4	3.959	-14.88
Node 22.23.3	3.5.8	18.98	8.447	4.021	6.515
Node 24.21.0	3.5.8	19.54	8.447	4.136	6.952
Node 26.10.0	3.5.8	18.01	8.447	3.421	6.145

Node 18.20.8 and 103% under Node 23.11.1, the excess over 100% being the difference between our build of the release and Node’s. There the discarded attempt *is* the OpenSSL parse, and it would cost the same with no JavaScript engine present. Second, the exception is small and nearly constant across runtimes: the Node-shaped error costs 4.781 μs under Node 18.20.8 and 3.421 μs under Node 26.10.0 (1.2% and 19.0% of the respective probes), and roughly half of it is stack capture (with capture disabled, 2.204 and 1.930 μs). Disabling stack capture for the real probe saves 5.397 μs and 1.767 μs on the same two runtimes, the same order as the synthetic figure. Third, on runtimes bundling OpenSSL 3.5, where the parse is cheap, the three parts are of comparable size: under Node 26.10.0 the parse is 47% of the probe, the exception 19.0% and the binding 34.1%.

The component is therefore a discarded *parse attempt*, not an exception cost. Where it is expensive the exception is about a hundredth of it, and even where it is cheap the exception is a fifth. These runtimes are Linux builds on x86_64; the ARM64 host figure of Table 2 was not itself decomposed this way, and its 18.04 μs is of the same order as the OpenSSL 3.5 rows here, not the OpenSSL 3.0 rows.

560 *4.3. The parts do not account for all of it*

561 If the penalty for supplying a string were exactly the discarded attempt
 562 plus the conversion, the two would sum to it. They do not, quite:

	observed penalty	23.92 μs
563	discarded parse + conversion	18.62 μs
	residual	5.292 μs

564 All three quantities are medians of per-round paired differences, per
 565 Section 3.2. The residual’s interval, [4.916, 5.501], excludes zero. The two
 566 measured parts therefore account for 77.9% and not more; the remainder
 567 lies in the library’s own type and claim checking along the non-pre-parsed
 568 path. We report it rather than absorbing it into the headline, because a
 569 decomposition whose parts are declared to sum exactly to the whole has
 570 usually not been checked.

571 The pre-parsed path carries a smaller unexplained remainder of the same
 572 kind, and we record it for consistency rather than leaving the asymmetry
 573 unremarked: `verify()` with a pre-parsed key costs 6.042 μs against 1.292 μs of
 574 structural decode plus 2.166 μs of HMAC, leaving roughly 2.58 μs unattributed.
 575 Neither residual was measured directly; both are the subject of conditions we
 576 did not run, and both are stated as unattributed rather than assigned.

577 *4.4. Confirmation without differencing*

578 The decomposition reaches its conclusion by subtracting conditions, the
 579 same inference style this paper criticises elsewhere. It is therefore checked
 580 against V8’s sampling profiler, which attributes self time to named frames
 581 with no arithmetic on our part. At 50000 iterations per run and 3 runs per
 582 environment, matched so the percentages are comparable, the asymmetric key
 583 parse accounts for a median 66.2% of self time on the host (range 65.3–67.4
 584 across runs), 95.3% (95.1–95.6) under Node 18, and 61.5% (58.5–61.6) under
 585 Node 26. The frame does not appear among the top frames of the pre-parsed
 586 path in any environment.

587 Two methods, one answer. We rely on the ordering and the dominance
 588 rather than on the exact percentages: sampling attribution depends on rate
 589 and on what else the process is doing.




Figure 4: The four container rows of Table 4 on a logarithmic axis. The discarded parse spans $25.36\times$ across images while the pre-parsed call spans $1.26\times$. The dashed rule marks the OpenSSL series boundary. Within this matrix it is not a controlled boundary, because the two fast images also carry a newer engine and a different build (Section 5.2); Section 5.5 controls it. The uncontainerised host row is omitted because its comparison with the container rows is confounded.

5. RQ2: What governs the magnitude

Table 4 varies the runtime with containerisation held constant, and Figure 4 plots the two columns that matter.

Table 4: The same measurement across four container runtimes on one machine, 10 processes each of 20000 calls. The *Disc. parse*: the discarded asymmetric parse; *Pre-parsed*: `verify()` with a pre-parsed key. [‡]The host row is from Table 2 and is not containerised. [†]Inferred from the run timestamp, not recorded per row: that run predates the per-measurement OpenSSL field (Section 10).

Runtime	OpenSSL	V8	Disc. parse (μs)	Pre-parsed (μs)
Node 18 musl	3.0.16	10	373.79	7.58
Node 20 musl	3.0.19	11	420.62	7.25
Node 26 musl	3.5.7	14	22.94	6.56
Node 26 glibc	3.5.7	14	16.58	6.04
host [‡]	3.6.3 [†]	14	18.04	6.042

5.1. Engine major version does not distinguish the two slow runtimes

Node 18 and Node 20 ship V8 10 and 11 (different major versions of the engine) and share the OpenSSL 3.0 series. Both pay of order $373.79\text{--}420.62\ \mu s$. Node 26 ships OpenSSL 3.5 and pays $22.94\ \mu s$ on the same libc. Within the

597 OpenSSL 3.0 series, then, changing the engine’s major version does not move
598 the quantity; had Node 20 been fast, the attribution would have been to the
599 engine.

600 *5.2. What this design cannot separate*

601 That argument establishes a negative within one OpenSSL series. It does
602 not establish the positive attribution for the change between series, and we
603 state the limitation rather than leaving it implicit.

604 The fast rows are Node 26. They carry OpenSSL 3.5 *and* V8 14 *and* every
605 other change in that build. OpenSSL version is therefore perfectly confounded
606 with engine version and with build configuration across the 3.0–3.5 boundary.
607 Nothing in this matrix can separate them. The absence of variation between
608 V8 10 and V8 11 does not license the inference that V8 14 is also irrelevant;
609 that is a different comparison, and we did not make it.

610 What the matrix supports on its own is the following, and no more: the
611 magnitude co-varies with the bundled OpenSSL version across four images;
612 the two slow images differ in engine major version and do not differ in
613 magnitude; and the two libc variants of one image differ by less than the
614 difference between series. Separating OpenSSL from the engine requires a
615 crossed design: a runtime pairing a newer engine with an older OpenSSL, or
616 the parse measured outside Node entirely. Section 5.5 reports both, and it is
617 on them, not on this matrix, that the causal wording used in this paper rests.

618 *5.3. The effect is confined to key parsing*

619 Across the four container rows, pre-parsed validation spans $1.26\times$, the
620 HMAC primitive $1.39\times$ and structural decode $1.35\times$. The discarded parse
621 spans $25.36\times$. The successful parse moves with it ($11.74\times$), so this is the cost
622 of key parsing in general rather than anything peculiar to the failure path.
623 The failure is simply the case in which an application pays that cost on every
624 request for a result it discards.

625 *5.4. What this design does not isolate*

626 A reader may notice that the uncontainerised host row is of the same
627 order as the Node 26 container row. That comparison is confounded and we
628 draw no conclusion about containerisation from it: the host runs a different
629 operating system from the containers, so operating system, C library and Node
630 build all change alongside the container boundary. Isolating containerisation
631 would require a Linux-native baseline, which we did not obtain. The claims




Figure 5: The failed parse of an HMAC secret issued from C, with no JavaScript engine, against 8 OpenSSL releases built from source with one configuration (Table 5, column *Full*), beside a successful parse of an RSA-2048 public key through the same path. Whiskers are 95% percentile bootstrap intervals over 15 processes and are narrower than the markers. The cost falls in two steps, at 3.1 and 3.2, and is flat thereafter.

above concern only the four container rows, among which containerisation is constant.

5.5. Caused by OpenSSL: the parse without an engine

Table 5 removes Node entirely: the OpenSSL calls `createPublicKey()` makes for an HMAC secret, issued from C against 8 OpenSSL releases built identically, and Figure 5 plots its first and last columns.

In plain C the failure costs 394.4 μ s on OpenSSL 3.0.16 and 8.447 μ s on 3.5.8, a ratio of 46.69 [46.16, 46.90]. Nothing else differs between those two measurements: the program, the compiler, the flags, the machine and the protocol are the same. The magnitude is therefore *caused* by the OpenSSL release. That is a claim about this code path in these releases; it says nothing about OpenSSL performance generally.

The table also says where in OpenSSL the cost lives and when it went away. Almost all of it is the private-key parse (99% of the full sequence on 3.0.16), the step at which OpenSSL 3 hands the input to its provider-based decoders; the three public-key scans reject the input at the PEM header and are cheap in every release, and the error strings cost under a microsecond. The fall came in two steps: to 57.48 μ s in 3.1.8, and to 6.998 μ s in 3.2.6, after which it is flat to within a few microseconds through 3.6.1. We measured the

Table 5: The failed asymmetric parse issued from C, with no JavaScript engine, against 8 OpenSSL releases built from source with one compiler and one set of flags (x86_64), in μ s. Medians of per-process medians over 15 processes of 5000 calls. *Full*: everything `createPublicKey()` asks of OpenSSL for an HMAC secret, including the error-string drain; 95% intervals for it are in the replication package and are within 0.9% of the median for every release. Its parts follow. *Successful parse*: an RSA-2048 public key through the same path, for comparison.

OpenSSL	Full	Private-key parse	Public-key scans	Error strings	Successful parse
3.0.16	394.4	390.9	2.287	0.796	112.9
3.0.19	431.3	424.0	2.231	0.869	132.6
3.1.8	57.48	54.84	2.196	0.497	36.09
3.2.6	6.998	4.487	2.057	0.459	11.75
3.3.5	6.946	4.465	2.074	0.474	12.04
3.4.3	7.245	4.554	2.098	0.500	11.19
3.5.8	8.447	5.757	2.051	0.483	11.26
3.6.1	8.803	5.985	2.234	0.472	11.16

steps, not the source changes that produced them; identifying those would take a bisection of the OpenSSL history that we have not done. A successful parse of real key material fell with it, which is why the effect is confined to key parsing in general rather than to the failure path.

Engine and library crossed. The official runtimes of Section 3.6 contain the pairing stock container images lack: Node 22.23.3 ships V8 12.4.254.21 with OpenSSL 3.5.8, while Node 23.11.1 ships the *newer* V8 12.9.202.28 with the *older* OpenSSL 3.0.16. The probe costs 18.98 μ s on the first and 383.5 μ s on the second. The Node 18/20 pair of Table 4 showed that a change of engine at a fixed OpenSSL series leaves the cost where it was; this pair shows that a newer engine does not rescue an older OpenSSL. Together with the C measurement, the engine is excluded from both directions.

5.6. A second architecture

Table 6 repeats the runtime matrix, with the same images and scripts, on x86_64 (Section 3.8). Absolute values are not compared across the two machines; the question is whether the ordering holds.

The property the attribution rests on (every OpenSSL 3.0 image slower than every OpenSSL 3.5 image) holds on both architectures. The strict rank

Table 6: The runtime matrix of Table 4 on two architectures, same images and scripts, in μs . ARM64: 10 processes of 20000 calls. x86_64, a shared cloud virtual machine: 10 processes of 20000 calls. Absolute values are not comparable across the two machines.

Image	OpenSSL	Discarded parse		Pre-parsed <code>verify()</code>	
		ARM64	x86_64	ARM64	x86_64
<code>node:18.20.8-alpine</code>	3.0.16	373.79	607.1	7.58	12.66
<code>node:20-alpine</code>	3.0.19	420.62	556.3	7.25	11.89
<code>node:26.6.0-alpine</code>	3.5.7	22.94	42.13	6.56	10.94
<code>node:26.6.0-bookworm</code>	3.5.7	16.58	19.91	6.04	9.701

order does not: the two OpenSSL 3.0 images, Node 18 and Node 20, exchange places between ARM64 and x86_64. Their difference is small against the gap between series on either machine, and no claim in this paper depends on their relative order. On x86_64 the discarded parse spans $30.50\times$ across the images while pre-parsed validation spans $1.31\times$, against $25.36\times$ and $1.26\times$ on ARM64. The host decomposition of Table 2 was repeated too (Node 22.22.2, OpenSSL 3.5.5): the discarded parse is $19.42 \mu s$ of a $36.48 \mu s$ call (53%), the symmetric constructor $1.261 \mu s$ and the HMAC primitive $3.105 \mu s$, so the component ordering of Table 2 holds there as well. The full x86_64 host table is in the replication package.

6. RQ3: The containerised service, and prevalence

6.1. In-situ cost across arrival rates

The containerised service was driven at seven arrival rates, each measured by the paired-scrape difference of Section 3.9, with the whole sweep run ascending and then descending on one process. Table 7 reports the range across the two passes at each rate for the configuration the service deploys, a string secret, and Figure 6 plots it.

One observation follows, and it is the one this section claims: the cost is present in the containerised service and is far larger than the isolated host figure of Table 2, consistent with Section 5 in that the service’s runtime is the one carrying the older OpenSSL.

Two things visible in Table 7 need care, and we set them out explicitly.

Wall-clock exceeds the prediction; processor time does not. The service runtime’s discarded parse costs $373.79 \mu s$ (Table 4), and the remainder of the




Figure 6: In-situ validation cost against arrival rate, string secret. The two passes are drawn separately rather than averaged: Section 6.2 reports that they differ systematically, and averaging would conceal that. The shaded band is the range between the passes, which is two observations rather than an interval.

Table 7: In-situ validation cost with a string secret, by arrival rate. The two columns are the ascending and the descending pass over the same seven rates on one process; each is a single window, and the pair is two observations rather than an interval. Host load average during each window is committed with the measurement.

Arrival rate (requests/s)	Mean per call (μ s)	
	ascending	descending
5	928.24	1027.97
10	976.28	1017.34
25	918.54	934.47
50	758.57	737.12
100	571.59	607.78
200	494.91	532.50
400	452.82	467.62

validation path is of order 7.58μ s, giving an expected total near 381μ s. The wall-clock means observed across the sweep are $453\text{--}1028 \mu$ s, exceeding that by $71\text{--}647 \mu$ s. Contention inflates a duration rather than deflating it, so these are upper bounds, and Section 6.4 shows where the excess is not: measured as processor time, the validation call costs within 0.8% of the prediction. The excess is therefore time the call spent off the processor on a shared host, not additional work. We did not decompose that off-processor time further, and no claim below depends on it.

The decline with arrival rate is not interpreted. Cost falls monotonically from 928.24μ s at 5 requests/s to 452.82μ s at 400 requests/s. This is the opposite of what contention alone would produce, and several mechanisms unrelated to validation would produce it: a fixed per-window overhead amortising over more calls, processor frequency scaling and idle-state exit latency at low rates, cache and branch-predictor residency, or continued tier-up in the engine. We did not run the conditions that would distinguish them. We therefore report the numbers and do not claim that validation cost depends on arrival rate; establishing that would require varying window duration at fixed rate, pinning the frequency governor, and extending warmup at low rates, none of which we did.

712 6.2. *The two passes differ systematically*

713 Sweeping up and then down is not only replication; it detects warm-up,
714 thermal and hysteresis effects. If the difference between passes were noise,
715 its sign would alternate across rate points. It does not. In the string-secret
716 sweep the descending pass is slower at 6 of 7 rate points; in the pre-parsed
717 sweep the descending pass is *faster* at 7 of 7. A consistent sign across seven
718 points is not what noise produces.

719 This has a direct consequence for how the ranges in Table 7 should be read.
720 A systematic component means the two passes are not independent samples of
721 a common value, so the range between them understates uncertainty further
722 than a range of two already does. We therefore report ranges and draw no
723 inference that depends on their width. It is also a methodological finding in
724 its own right: a single-direction sweep would have produced a curve carrying
725 a consistent bias with no way to detect it.

726 6.3. *Host load, not the condition, drives the outliers*

727 Agreement between passes varies considerably across windows. We dis-
728 tinguish *relative* disagreement, the gap between passes as a fraction of their
729 mean, from *absolute* disagreement, the gap in microseconds; the two behave
730 differently and the distinction carries this subsection.

731 The widest relative disagreement in the dataset and the tightest occur at
732 the same arrival rate in different sweeps. The wide one is the highest-load
733 window in its own sweep on the descending pass, though not on the ascending
734 pass, whose maximum load falls at a different rate. The tight one, at the
735 same arrival rate in another sweep, sits at a lower load on both passes and
736 its two passes agree to within $0.02\ \mu\text{s}$. Relative disagreement therefore does
737 not track the condition under measurement.

738 We considered two alternative explanations for the relative asymmetry and
739 reject both. A small-mean effect would require absolute disagreement to be
740 roughly constant, so that its relative size grew only as means shrank; absolute
741 disagreement instead varies over three orders of magnitude and is largest
742 where the means are largest, which is the opposite pattern. A measurement
743 floor would require the affected values to sit at the instrument’s resolution;
744 they do not, and in any case the means reported here are computed from the
745 histogram’s sum and count rather than from its buckets, so bucket resolution
746 cannot bound them.

747 The operative consequence is that in-situ figures are reported with the
 748 host load recorded during their window, and comparisons between windows
 749 at different loads are not treated as comparisons between conditions.

750 *6.4. Corroboration: the cost is processor time, and it matches the prediction*

751 An in-situ figure of this magnitude has two readings. Either it is work the
 752 service performs, or it is time the request spends waiting, queued behind other
 753 requests, descheduled by a contended host, or inflated by the instrument.
 754 We separate them with an A/B of the same service at one arrival rate, with
 755 validation enabled in one window and disabled in the next, on the same
 756 pod running the same image, and read the processor time the application
 757 container was charged from its cgroup accounting rather than from a clock.
 758 Table 8 reports it.

Table 8: Validation enabled against validation disabled, at one arrival rate, two 180 s windows on the same pod and image, 36001 requests each. A signed token is on the wire in both arms; the disabled arm removes only the handler’s `jwt.verify()` call. Application and sidecar CPU are means of one-second cgroup samples in the window; CPU per call is that mean divided by the achieved rate. Event-loop lag is the mean of the per-scrape p99. Host load is a single sample at window start.

	disabled	enabled
Achieved rate (req/s)	200.006	200.006
Application CPU (millicores)	69.45	146.38
CPU per call (μs)	347.3	731.9
Sidecar CPU (millicores)	53.40	49.13
Event-loop lag (ms)	3.137	3.159
Client mean latency (ms)	1.6802	2.0332
Host load, 1 min	3.53	5.11

759 Enabling validation adds 384.6 μ s of processor time per call. The runtime
 760 matrix predicts 381 μ s for the same work on the same runtime (Section 6.1);
 761 the two agree to within 0.8%. They are obtained by instruments that share
 762 no code: cgroup CPU accounting on the service’s pod, and a single-process
 763 microbenchmark in an isolated container. A request that queues or is de-
 764 scheduled accrues wall-clock time but no CPU, so this figure cannot contain
 765 queueing delay.

766 Three observations weigh against reading the difference as an artefact of
 767 the two windows. The sidecar, in the same pod under the same contention, did

not rise with the application: 53.40 millicores disabled against 49.13 enabled. The one-second CPU samples of the two arms do not overlap at all (disabled never exceeded 100.577 millicores in 127 samples, enabled never fell below 114.963 in 127), which separates the arms without a distributional assumption. And event-loop lag is unchanged, so neither arm was near saturation. The two windows are two observations, not a sample: the difference carries no interval and we compute none.

6.5. *How common is the avoiding pattern, in our sample?*

Three instruments, three frames, reported separately as Section 3.10 requires. All counts were captured on 2026-09-18. None of the figures below is a population proportion.

Population co-occurrence. Of 790528 public JavaScript files importing the library through `require`, 29 also mention the pre-parsing interface; for JavaScript through `import` the figures are 524288 and 38; for TypeScript through `require`, 8288 and 0; and for TypeScript through `import`, 333824 and 101. Cells are not pooled into a rate: a file using both import forms would be counted twice, so their sum 1656928 is an upper bound on the union and the largest cell 790528 a lower bound. No cell exceeds 0.0303% of files mentioning the interface. These are search-engine estimates of file counts over public repositories only, they measure co-occurrence rather than use, and the denominators are returned *rounded* (each is an exact multiple of 1024), so they should be read as magnitudes rather than as counts established file by file.

Classified call sites. Of 327 call sites classified across 247 distinct repositories in our sample, 0 pass a pre-parsed key. Environment variables account for 176 of those call sites, untraced identifiers and member expressions 93, string literals 52, file contents 4 and buffers 2. No file in this corpus contains the pre-parsing interface at all, in any spelling, so the zero reflects the pattern's absence from the sample rather than a classifier that cannot see it. Incidentally, 309 of 327 call sites (94%) pass no algorithm restriction, which Section 8.4 returns to.

Matching only the conventional `jwt.verify()` spelling would filter the corpus toward one import idiom, so call sites are found by resolving the local binding of the import, covering arbitrary aliases, namespace imports, destructured and renamed imports, and direct calls on the `require` expression. That

803 binding resolution contributes 12 of the call sites above that the conventional
804 spelling misses, of which 0 pass a pre-parsed key.

805 *Targeted counter-search.* A sample can only show a pattern is rare if it
806 could have found it. Every file co-occurring the library with the pre-parsing
807 interface was therefore fetched: 191 files, of which 120 contain a validation
808 call, spanning 112 distinct repositories. An automated pass flagged 13 call
809 sites as passing a pre-parsed key and all were then read by hand: 5 are false
810 positives (projects defining their own helper function of the same name, one
811 of which returns a string), 1 could not be resolved within its file, and 7 are
812 genuine, across 6 distinct projects.

813 *The rate, with its denominator.* Among the 112 repositories that both refer-
814 ence the pre-parsing interface and call the validator, 6 use it: 5.36%. That
815 denominator is the most favourable available, since those projects demonstra-
816 bly know the interface exists, and it is not the population. Within this frame
817 the pattern is rare rather than absent, and it is rare even among projects
818 already aware of it. We do not extrapolate beyond the frame.

819 7. RQ4: One library, or a design choice?

820 A cost found in one library is an observation about that library. It becomes
821 a statement about a class of defect only if the mechanism can be named
822 independently of the library and predicts which other libraries have it. The
823 mechanism here can be named: *resolving key material of unknown type by*
824 *attempting the asymmetric parse first, and treating its failure as the signal*
825 *that the material is a shared secret.* Table 9 tests the prediction against six
826 further libraries.

827 *What the source says.* Only `jsonwebtoken` attempts the asymmetric parse.
828 The others resolve the ambiguity in one of three ways, none of which involves
829 a parse that can fail on a shared secret. *Inspect, then dispatch:* `fast-jwt`
830 examines the text of the key (whether it carries a PEM header or is a JSON
831 Web Key) and calls exactly one constructor, skipping even that inspection
832 when every permitted algorithm is HMAC; `PyJWT` likewise inspects the text,
833 and rejects PEM- or SSH-looking material offered as an HMAC key rather
834 than trying to parse it. *Refuse the ambiguous input:* `jose` does not accept
835 a string for an HMAC key, and `golang-jwt` accepts only a byte slice for
836 HMAC verification, returning a key-type error for anything else. *Type the*

Table 9: HS256 verification across seven JWT libraries in four languages (x86_64), in μ s. *Asym. first?*: given a string HMAC secret, does the library attempt to parse it as an asymmetric key, as read from its source. *String* and *pre-parsed*: the secret in the form the library’s documentation shows, and in the library’s own key object built once (Section 3.7). *Penalty*: median of per-round paired differences, with a 95% bootstrap interval over rounds. Medians of per-process medians over 15 processes of 20000 calls. Absolute costs measure different amounts of work in different runtimes and are not a ranking of libraries. ^aNo string accepted for an HMAC key; the rawest accepted form (bytes) is used. ^bA string is accepted only by a deprecated setter that resolves it once, when the parser is built.

Library	Lang.	Version	Crypto backend	Asym. first?	String	Pre-parsed	Penalty	95% interval
jsonwebtoken	JS	9.0.3	OpenSSL 3.0.19	yes	476.2	10.23	465.9	[465.2, 467.8]
jsonwebtoken	JS	9.0.3	OpenSSL 3.5.8	yes	34.22	9.844	24.56	[24.07, 24.82]
fast-jwt	JS	6.3.3	OpenSSL 3.0.19	no	7.564	7.570	0.070	[-0.116, 0.111]
fast-jwt	JS	6.3.3	OpenSSL 3.5.8	no	7.917	7.934	-0.068	[-0.134, 0.197]
jose	JS	6.2.12	OpenSSL 3.0.19	no ^a	121.2	78.42	42.66	[38.50, 45.89]
jose	JS	6.2.12	OpenSSL 3.5.8	no ^a	98.89	73.98	24.80	[21.91, 26.51]
PyJWT	Python	2.15.0	hashlib, OpenSSL 3.0.13	no	41.33	41.28	-0.103	[-0.457, 0.338]
golang-jwt	Go	5.3.1	Go crypto/hmac	no ^a	4.810	4.764	0.034	[-0.042, 0.117]
nimbus-jose-jwt	Java	10.10	JCA SunJCE	no	4.358	4.229	0.304	[-0.098, 1.016]
jjwt	Java	0.13.0	JCA SunJCE	no ^b	12.15	12.33	-0.406	[-0.715, 0.648]

837 *key*: jjwt requires a **SecretKey** for HMAC and a **PublicKey** for signatures,
838 so the distinction is made by the Java type system; nimbus-jose-jwt offers
839 a constructor that takes a string and converts its bytes without parsing them.
840 Refusing a string does not by itself remove per-call work, however: jose,
841 handed raw bytes, imports them into a WebCrypto key on every call.

842 *What the measurements say*. Of the seven libraries, two pay for being handed
843 their raw form rather than a key object, in the sense that the paired interval
844 on the penalty excludes zero; for the other five it does not, and their point
845 estimates are within a microsecond of zero. The two are jsonwebtoken
846 and jose, and they pay for different reasons. jsonwebtoken performs the
847 failing asymmetric parse: its penalty is 465.9 μ s under OpenSSL 3.0.19 and
848 24.56 μ s under OpenSSL 3.5.8, a string-to-pre-parsed ratio of 46.55 and 3.48
849 respectively. jose parses nothing that can fail, but imports the raw bytes
850 into a WebCrypto key on every call: 42.66 μ s and 24.80 μ s under the same
851 two runtimes. Both penalties grow on OpenSSL 3.0, so the OpenSSL release
852 sets the price of per-call key handling generally; but the failing parse is the
853 costliest form of it, about 11 times jose’s on OpenSSL 3.0. Two of the
854 libraries that pay nothing resolve the key once, when a verifier or parser is
855 built (fast-jwt, jjwt), so a raw secret costs nothing per call by construction;
856 that is itself the design point.

857 *What this establishes.* The libraries without the problem matter as much as
858 the ones with it. All seven validate the same kind of token with the same
859 primitive (against the same OpenSSL, in the JavaScript cases) and five pay
860 nothing for a raw secret. The cost therefore follows a design choice, resolving
861 key material on every call, and not JSON Web Token validation, HMAC or
862 OpenSSL as such: OpenSSL 3.0 sets the price, but only a library that makes
863 that choice is charged it, and one that makes it by way of a failing parse is
864 charged the most. The failing parse is exactly what the order-of-evaluations
865 profiler of Selakovic et al. (2017) is designed to find, and the alternatives
866 above are the fixes that line of work suggests: reorder the test, make it cheap,
867 or move it off the per-call path. The comparison does not estimate how many
868 libraries make the choice; it shows that making it is what matters.

869 **8. RQ5: The cost as a security property**

870 Sections 4–7 treat the discarded attempt as a cost. This section asks
871 who can cause it, why it is there, and what removing it does to the library’s
872 defences. Its evidence is the library’s source at the releases named, the
873 library’s own test suite, and the measurements already reported; it introduces
874 no new timing condition, and where a statement rests on reading code rather
875 than on measurement we say so.

876 *8.1. Threat model*

877 The adversary is a remote client that can place arbitrary bytes where the
878 service expects a bearer token, including every field of the token header, and
879 knows which library the service uses. It does not know the shared secret
880 and cannot produce a valid HMAC. The service verifies HS256 tokens with a
881 shared secret supplied as a string, the form the library’s documentation shows
882 and the form every call site in our sample uses (Section 6.5). The asset of
883 interest is key-type separation: a token declaring an HMAC algorithm must
884 never be verified against asymmetric key material, nor the reverse (Sheffer
885 et al., 2020). Out of scope are compromise of the secret, defects elsewhere
886 in the library, and availability. We did not drive adversarial traffic at the
887 containerised service and make no availability or capacity claim; the question
888 here is which inputs reach the discarded attempt, not how many a service
889 can absorb.




Figure 7: Order of operations in `verify()` at tag v9.0.3, with the file’s line numbers. Every step up to and including the algorithm and key-type checks runs before the signature is verified, and so runs for a forged token exactly as for a genuine one. The discarded attempt sits inside key resolution; the key-type checks that implement the library’s key-confusion defence read the type of the key object resolution produces (dashed arrow). Pinning permitted algorithms is enforced at the third step and therefore does not prevent the second.

8.2. The attempt precedes authentication

Figure 7 shows the order in which `verify()` performs its checks at tag v9.0.3. Reaching the key-resolution block of Figure 2 requires only that the token have three dot-separated segments, a header and payload that decode as JSON, and a non-empty signature segment (lines 59–118). Nothing on that path depends on the secret. The permitted-algorithm list is consulted afterwards, at line 144; the key-type checks at lines 148–160; the signature at line 165.

Two consequences follow. A token that will be rejected pays the discarded attempt in full: 18.04 μ s on the host, and 373.79 μ s on the runtime the studied service deploys (Table 4). And restricting permitted algorithms, the defence RFC 8725 recommends and which 94% of the call sites in our sample omit, does not avoid it, because the restriction is enforced after resolution. The cost is therefore under the control of any client that can reach the endpoint, not only of the service’s legitimate load. That is a statement about reachability, established from the source; its magnitude is the one measured in Sections 4–6.

8.3. The attempt arrived with a key-confusion fix

We compared `verify.js` in the published packages of versions 8.5.1, 9.0.0 and 9.0.3. In 8.5.1 the secret is passed to the signature routine as given, with no key-object resolution. The resolution block of Figure 2 first appears in 9.0.0

910 and is unchanged, apart from whitespace, through 9.0.3. Version 9.0.0 is the
911 release that fixed CVE-2022-23541 and CVE-2022-23539 (Auth0, 2022), and
912 the checks implementing those fixes (the comparison of the declared algorithm
913 with the key’s type at lines 148–152 and the validation of asymmetric key
914 types at lines 154–160) operate on the key object the resolution block produces.
915 The attempt is thus not an oversight in otherwise untouched code. It is part of
916 a security fix, and its price is set by whichever OpenSSL release a deployment
917 happens to bundle (Section 5).

918 This is consistent with the one earlier report of the symptom: HS256
919 verification in the 9.x series measured 34 times slower than in 8.5.1, on a
920 Node release bundling OpenSSL 3.0 (Section 8.6). We did not time version
921 8.5.1 ourselves, and state the connection as consistent with the report rather
922 than as established by it.

923 *8.4. A fix, and a constraint on it*

924 The fix is to dispatch on the token’s declared algorithm before attempting
925 the asymmetric parse, so that an HMAC token with an ordinary string secret
926 is resolved directly. Applied, this takes the call from 29.92 μ s to 7.208 μ s, a
927 factor of 4.15, and to within 1.208 μ s of supplying a pre-parsed key, with no
928 change in calling code. That residual is the median of the per-round paired
929 differences, not the difference of the two medians; the two are not identical
930 and the former is reported throughout.

931 We tested it against the library’s own suite rather than only against
932 assertions of our own. Cloned at tag v9.0.3, the patch applies cleanly and the
933 suite is unchanged: 511 passing and 0 failing unpatched, 511 and 0 patched.

934 These figures are from the host. The practical benefit is largest where the
935 discarded parse is largest, of order 373.79 μ s on the runtime the containerised
936 service uses (Table 4), and we did not measure the patched library on the
937 container runtimes. The reported factor of 4.15 is therefore a lower bound
938 on the improvement available in the containerised service’s runtime, and we
939 state it as a host result rather than as a general one.

940 *The constraint, which is the more important half.* The discarded attempt is
941 not purely wasteful. The *success* of that parse on asymmetric key material
942 is what produces a key object whose type a later check depends on, and
943 that check is what prevents symmetric and asymmetric material from being
944 interchanged (Sheffer et al., 2020). A change that dispatches on the declared

945 algorithm while relaxing which key material may take the fast path alters
946 which inputs reach that check.

947 The fix must therefore be narrow: the fast path is restricted to string
948 secrets that cannot be asymmetric key material, so anything that might
949 be resolved exactly as it does today. Removing that material restriction
950 (the form of the optimisation we argue against) produces 2 failures in the
951 library’s own suite, both in its key-confusion tests (“should not allow HMAC
952 verification with an RSA key in PEM format” and, under “signing with pub
953 key as symmetric”, “should not verify”). All of these outcomes (the suite
954 unchanged under the narrow patch, two key-confusion failures without its
955 guard) reproduce on the x86_64 machine of Section 3.8 under a different Node
956 and OpenSSL release (Katha et al., 2026).

957 *The existing suite already guards the property.* No new regression test is
958 needed: the library’s existing suite constrains the change, and a contributor
959 proposing the unrestricted optimisation would be caught by tests they are
960 required to run. The finding is therefore not a missing test but a hidden
961 coupling: the cost documented in Section 4 cannot be removed by the obvious
962 route, and the reason is a property the maintainers already test for but which
963 is not evident from the code being changed.

964 The independent defence against this class of confusion is for the caller to
965 restrict permitted algorithms (Sheffer et al., 2020), and 94% of the call sites
966 surveyed in Section 6.5 do not.

967 8.5. *What the narrow fix leaves reachable*

968 The narrow fix selects its fast path on the token’s declared algorithm,
969 and the declared algorithm is chosen by whoever constructs the token. For a
970 token declaring an HMAC algorithm and a plain string secret the parse is
971 skipped. For a token declaring any other algorithm the patch falls through
972 to the original resolution, so `createPublicKey()` is attempted on the shared
973 secret and fails exactly as it does today; the token is then rejected by the
974 algorithm list or the key-type check, but only after the parse. Table 10 sets
975 out which inputs reach the attempt under each configuration. The narrow
976 fix removes the cost from legitimate traffic, not from traffic an adversary
977 constructs. Only a pre-parsed key object removes it for every input, because
978 the resolution block is skipped entirely when the configured key is already a
979 `KeyObject` (line 120).

Table 10: Which tokens reach the discarded key-parse attempt, for a service configured with an HMAC shared secret, as read from the source at tag v9.0.3 and from the patch of Section 8.4. “Well-formed”: three segments, JSON header and payload, non-empty signature. No row requires knowledge of the secret.

Configuration	Tokens that reach the parse
Stock, string secret	every well-formed token, forged or genuine
Narrow fix, plain string secret (no PEM or JWK form)	every well-formed token whose header declares a non-HMAC algorithm; no such token can be valid for this service
Pre-parsed <code>KeyObject</code>	none

980 This follows from the code and was not separately timed: we did not
 981 measure the patched library on tokens declaring a non-HMAC algorithm,
 982 and the per-call cost such a token incurs is that of the discarded attempt
 983 measured in Sections 4 and 5. A variant that dispatched on the *configured*
 984 algorithms rather than the declared one would close the gap for callers who
 985 restrict permitted algorithms to HMAC, since tokens declaring anything else
 986 are rejected regardless; it is the choice `fast-jwt` makes (Section 7). We did
 987 not implement or test it, it would help only the minority of callers who restrict
 988 algorithms, and it would have to keep the same guard on key material.

989 8.6. Reported upstream

990 Before reporting, we searched the library’s issue tracker for earlier reports.
 991 One exists: an issue opened in April 2024¹ reports that HS256 verification
 992 in the 9.x series is 34 times slower than in 8.5.1, on a Node release that
 993 bundles OpenSSL 3.0. It diagnoses no cause and, at the time of writing,
 994 has no maintainer response. That is the symptom this paper explains,
 995 observed independently two years earlier by a user who could not see its
 996 source. We reported the mechanism, with a minimal reproduction, per-
 997 runtime measurements and the OpenSSL finding, as a separate issue², and a
 998 pull request implementing the narrow form of the fix, with its guard, is open
 999 against it. At the time of writing neither has received a maintainer response.

¹<https://github.com/auth0/node-jsonwebtoken/issues/966>

²<https://github.com/auth0/node-jsonwebtoken/issues/1046>

9. Implications for practice

Pre-parse key material once, at startup. On the host this removes 23.92 μs per validation [23.58, 24.00]; on the runtime the service actually deploys, where the discarded parse costs 373.79 μs , it removes far more. It requires no architectural change, no new component and no movement of a trust boundary, and it is the only remedy in this paper that removes the attempt for every input, including tokens an adversary constructs (Table 10). On a current runtime bundling OpenSSL 3.5 the saving is smaller (24.56 μs per validation in Table 9), but it is not negligible beside the 347.3 μs of processor time the containerised service spends on a request with validation disabled (Table 8). The two figures come from different machines and are compared as orders of magnitude only.

Treat the base image as part of the measurement. The same code, the same library and the same architecture differ by $25.36\times$ on this path across the container images measured, because they pin different OpenSSL versions. A performance figure quoted without the base image it was obtained on is not reproducible, and a regression appearing after a base-image bump may have nothing to do with the application. Section 5.5 shows that on this path the OpenSSL release is the cause; but the practical point would hold even if it were not, since it is enough that the image changes the number.

For library authors: do not use a failing parse as a type test. Section 7 found three designs that avoid the cost (inspect the material and dispatch, refuse ambiguous input at the interface, or type the key) and all three ship in maintained libraries today. What they share is that resolution happens once, or not at all, rather than on every call. The design that pays most is the one that asks an expensive parser to answer a cheap question on every request.

Pin the algorithm. Independently of anything here, callers should restrict permitted algorithms (Sheffer et al., 2020). Most of the call sites we surveyed do not, which leaves them relying entirely on library-internal key typing for a property they could assert directly in one argument. Pinning does not remove the cost measured here, which is incurred before the restriction is checked (Section 8.2); it is an independent defence, not a substitute for pre-parsing.

For security reviewers: review performance changes to verification as security changes. The cheapest-looking optimisation of this path fails the library’s

key-confusion tests, and the property it breaks is not evident from the lines being changed (Section 8.4). Conversely, the security fix that introduced the path carries a cost that varies by more than an order of magnitude with the bundled OpenSSL release (Sections 5 and 8.3). Changes to token verification should be run against the library’s key-confusion tests and timed on the oldest cryptographic library a deployment may bundle.

9.1. Relocation is not the first remedy

A cost believed to be cryptographic invites relocation to a sidecar or gateway. The cost measured here is not cryptographic and is largely avoidable in place, so an architectural remedy that also moves a trust boundary should not be the first thing reached for. We implemented that relocation for the service studied here; applying a mesh validation policy duplicates verification rather than moving it, and making it effective requires the application to trust a proxy-supplied claims header, a known anti-pattern that needs a filter stripping client-supplied copies. We report that only as a design note: it carries no measurement, and the configuration, the observation and the inconclusive automated probe are documented in the replication package (Katha et al., 2026).

10. Threats to validity and future work

The limitations below are grouped by what they bear on: the measurements, the deployment evidence, the survey of public code, and the scope of the claims. Each states what the paper does not claim as a result.

10.1. Measurement

Two architectures, one of them shared. The primary measurements come from a single ARM64 host. Section 5.6 repeats them on x86_64, and there the split between OpenSSL series and the component ordering of the host decomposition hold. But the x86_64 machine is a virtual machine on shared cloud infrastructure, whose absolute timings are known to be unreliable (Laaber et al., 2019); we therefore take only orderings and ratios between interleaved conditions from it, and no magnitude in this paper should be read as holding on x86_64 hardware in general. A repetition on a dedicated x86_64 host, which the replication package scripts in one command, would turn the ordering check into a replication of magnitudes. The C probe, the exception baseline and the cross-library comparison were measured only on the x86_64 machine.

1068 Their conclusions rest on ratios between conditions measured in the same
1069 rounds, which is the kind of quantity a shared machine can support, and
1070 the effects they report are far larger than the machine’s noise: the median
1071 between-process coefficient of variation is 1.5%, 2.9% and 2.8% in the three
1072 runs. The largest values (28.0%, 27.7% and 20.2%) occur on sub-microsecond
1073 or JIT-dominated conditions that carry no claim, and medians of per-process
1074 medians are insensitive to them. The C probe needs no JavaScript runtime
1075 and the scripts rebuild every OpenSSL release from source, so repeating it on
1076 the ARM64 host or a dedicated x86_64 machine is one command.

1077 *The C probe reproduces Node’s calls, not Node’s build.* Section 5.5 issues the
1078 OpenSSL calls Node issues, as read from the Node source, against OpenSSL
1079 releases we built. Node builds its bundled OpenSSL with its own configuration,
1080 so the C figures approximate rather than reproduce the parse cost inside a
1081 given runtime, and the binding column of Table 3 inherits that approximation;
1082 it is slightly negative on two OpenSSL 3.0 runtimes for that reason. The causal
1083 claim does not inherit it: it compares releases under one build configuration
1084 with nothing else varying.

1085 *The exception baseline is a reconstruction.* Node constructs the error from
1086 C++ through V8’s API; our baseline constructs an error of the same shape,
1087 message and properties from JavaScript at a comparable stack depth. The
1088 two paths share V8’s error construction and stack capture, and the paired
1089 stack-capture measurement on the real probe agrees in order of magnitude
1090 with the synthetic one, but they are not the same code path. The ARM64
1091 host figure of Table 2 was not itself decomposed; Table 3 decomposes the
1092 same component on six other runtimes.

1093 *The headline host figure has no per-row OpenSSL version.* The host decom-
1094 position of Section 4 predates the per-measurement OpenSSL field. The
1095 version in force is known from the run timestamp and a note recorded at the
1096 time: OpenSSL 3.6.3, which a package upgrade later that day moved to 3.6.4
1097 under the same Node binary. We report it as inferred, marked as such in
1098 Table 4, rather than as recorded. Three measurements that do carry their
1099 version bound how much this matters. The same probe costs 18.01 μ s under
1100 Node 26.10.0 with OpenSSL 3.5.8, 19.42 μ s in the x86_64 host repetition with
1101 OpenSSL 3.5.5, and 16.58 μ s in the Node 26 container with OpenSSL 3.5.7:
1102 all of the same order as 18.04 μ s, and all from OpenSSL releases after the
1103 step of Table 5. The headline figure is therefore representative of a post-3.2

1104 OpenSSL, which is all the paper uses it for; it would be wrong by more than
1105 an order of magnitude only if the host had run an OpenSSL 3.0 or 3.1 release,
1106 which the record excludes.

1107 *10.2. Deployment measurement*

1108 *Ranges of two, with a systematic component.* The in-situ dispersion is the
1109 range across an ascending and a descending pass. It is not a confidence
1110 interval, and Section 6.2 shows the two passes differ systematically rather
1111 than randomly, so they are not independent samples of a common value and
1112 the range understates uncertainty further. No inference in this paper depends
1113 on the width of those ranges.

1114 *One in-situ process, two passes.* The rate sweep rests on one process sweeping
1115 seven rates twice. That is sufficient for the claim made (the cost is present
1116 in the containerised service and of this order) and insufficient for any claim
1117 about the shape of the cost-versus-rate relation, which Section 6.1 accordingly
1118 does not make. The magnitude claim does not rest on the sweep alone: the
1119 A/B of Section 6.4 is a separate pair of windows, measured with a separate
1120 instrument, and it agrees with the microbenchmark prediction on processor
1121 time. The wall-clock excess over that prediction is attributed to time off
1122 the processor and not decomposed further; the decline with arrival rate has
1123 several candidate explanations unrelated to validation that we did not test
1124 (Section 6.1).

1125 *Host load, and the confound that motivated recording it.* A throughput curve
1126 measured on a host shared with its own load generator moves with unrelated
1127 activity on that host, as Section 3.11 describes, which is why no capacity claim
1128 appears in this paper. Every in-situ window here carries the load average
1129 observed during it, and Section 6.3 shows that window-to-window agreement
1130 tracks that load rather than the condition. Contention inflates a duration
1131 rather than deflating it, so in-situ figures are upper bounds; but comparisons
1132 between windows at different loads are not comparisons between conditions,
1133 and we do not treat them as such.

1134 *The A/B is two windows at unequal load.* The two arms of Table 8 are
1135 consecutive windows rather than interleaved rounds, and their host load
1136 is known only at the instant each began. That bars comparing them on
1137 wall-clock time, and we do not. It does not bar the CPU comparison the
1138 section rests on: contention inflates elapsed time, whereas a container's own

1139 cgroup accounting charges it only the processor time it used. Attaching the
1140 per-second host sampler to a re-run of both arms, interleaved, would remove
1141 the limitation; we did not.

1142 *10.3. Survey of public code*

1143 *Sampling frames are not interchangeable.* The three prevalence instruments of
1144 Section 3.10 have different units and different selection mechanisms, and none
1145 supports a population proportion. The co-occurrence counts measure two
1146 strings appearing in one file, with a false-positive rate established at 5 of 13 on
1147 hand inspection. The classified sample has unknown inclusion probabilities, a
1148 limitation of repository search noted more generally by Kalliamvakou et al.
1149 (2014). The counter-search is selected on the outcome.

1150 *The matcher’s coverage.* Call sites are found by resolving the local binding
1151 of the import, which covers aliases, namespace imports, destructured and
1152 renamed imports and direct calls on the require expression. It does not resolve
1153 bindings across files, so a project that constructs a pre-parsed key in one file
1154 and validates in another is invisible to this analysis. This is the strongest
1155 threat to Section 6.5; within-file tracing found no such case, which is weak
1156 evidence against it.

1157 *One rater.* The 13 automatically flagged call sites were read and classified
1158 by the first author, and the authors are not disinterested: the paper’s claim
1159 is stronger the more of them are false positives, and 5 were classified that
1160 way. There was no second rater and no inter-rater agreement to report. The
1161 per-case verdicts, the evidence for each and the hashes of the files they were
1162 read from are published so the adjudication can be repeated, which is weaker
1163 than having repeated it.

1164 *10.4. Scope*

1165 *Token caching is not in the comparison.* A service that caches verified tokens
1166 pays this cost once per token rather than once per request. We did not
1167 measure such a configuration, so the per-request figures here do not describe
1168 a caching deployment.

1169 *The security analysis is from source.* Section 8 establishes reachability and
1170 provenance by reading `verify.js` at the releases named and by running the
1171 library’s own tests; it does not drive forged tokens at the containerised service,
1172 time the patched library on tokens declaring a non-HMAC algorithm, or

1173 time version 8.5.1. Its statements are therefore about which inputs reach the
1174 attempt and when the code arrived, not about throughput under attack, and
1175 the paper makes no availability claim.

1176 *Seven libraries, one version each.* Section 7 covers seven libraries in four
1177 languages, chosen as the most used in each ecosystem, each at one release.
1178 Other libraries, and other releases of these, may resolve keys differently; each
1179 classification is from source at the version measured. The comparison supports
1180 the claim that the cost follows a design choice rather than JWT validation; it
1181 does not estimate how many libraries make that choice.

1182 10.5. Future work

1183 Four directions follow from the limitations above.

1184 *Security.* Section 8 establishes which inputs reach the discarded attempt; it
1185 does not measure what an adversary can do with it. Driving forged tokens at
1186 the containerised service, under the stock library and under the narrow fix,
1187 would turn that reachability statement into an availability measurement. The
1188 configured-algorithm variant of the fix (Section 8.5) should be implemented
1189 and run against the library’s key-confusion tests, and a fix that removes the
1190 attempt for every input without requiring callers to pre-parse keys remains
1191 to be designed and proposed upstream.

1192 *Cause.* The C probe locates the fall in two steps, at 3.1 and 3.2 (Section 5.5).
1193 Bisecting the OpenSSL history between those releases would identify the
1194 changes responsible, and timing version 8.5.1 of the library would test the
1195 connection to the earlier upstream report directly (Section 8.3).

1196 *Generality.* Seven libraries at one version each show that the cost follows a
1197 design choice; they do not estimate how often the choice is made. A systematic
1198 survey of JWT and JOSE libraries, and of other authentication paths that
1199 resolve key material on every call, would. Resolving call sites across files and
1200 adding a second rater would strengthen the prevalence evidence of Section 6.5.

1201 *Deployment.* A dedicated x86_64 repetition, a Linux-native baseline for con-
1202 tainerisation, interleaved A/B windows with per-second host sampling, a
1203 controlled study of the decline with arrival rate, and a deployment that caches
1204 verified tokens would each close one limitation stated above. The scripts for
1205 the first are already in the replication package.

11. Conclusion

The per-request cost of token validation in this library is real, is worth attending to, and is not the signature: the primitive is $2.208\ \mu\text{s}$ of a $29.92\ \mu\text{s}$ call. It is not key conversion either, which costs $0.583\ \mu\text{s}$. It is a parse attempted in the wrong order, whose failure is constructed and thrown away on every request at $18.04\ \mu\text{s}$, 60% of the call. Where that attempt is expensive it is almost entirely OpenSSL work; the exception that ends it costs a few microseconds.

It is also a security property of the path. The attempt runs before the algorithm check and the signature, so any well-formed token can cause it; it entered the library with a key-confusion fix and that fix depends on it, so removing it naively fails the library's key-confusion tests; and the narrow fix that keeps the defence still leaves it reachable through an attacker-chosen algorithm header. Pre-parsing the key once, at startup, is both the fastest configuration and the only one no input can push onto the slow path.

Its price is set by the OpenSSL release a base image happens to pin, varying $25.36\times$ across runtimes while the rest of the validation path varies by $1.26\times$, and set causally, not only by association: the same calls from C cost $394.4\ \mu\text{s}$ on OpenSSL 3.0.16 and $8.447\ \mu\text{s}$ on 3.5.8, and a newer engine paired with an older OpenSSL is the slow runtime, not the fast one.

It is a design choice, not a property of JWT: of six other libraries in four languages none attempts the parse, five pay nothing for a raw secret, and the one that re-imports its key on every call pays an order of magnitude less. The symptom had been reported upstream two years before this study without a diagnosis; we have reported the mechanism.

The pattern that avoids the cost is rare in our sample: among 112 repositories that already reference the pre-parsing interface and call the validator, 6 use it. No population proportion follows.

The explanation that presents itself before the measurement, that the dominant cost is key conversion, is refuted by measuring the conversion. It has the shape this paper's methodology is built to avoid: a cost attributed to the variable under study while a second variable moves alongside it, unrecorded at the granularity at which it could change. The bundled OpenSSL version in Section 5 is the same hazard caught in time: the quantity varies $25.36\times$ with a component most measurement protocols record once per run, if at all. The lesson is not that one should measure in deployment. It is narrower: record the identity of the component whose behaviour is being measured, on every

1243 measurement, at the granularity at which it can change.

1244 **CRedit authorship contribution statement**

1245 **Jannatul Ferdous Katha:** Conceptualization, Methodology, Software,
1246 Investigation, Data curation, Formal analysis, Visualization, Writing – original
1247 draft. **Tasmia Tahmid Prova:** Validation, Writing – review & editing. **Md**
1248 **Kishor Morol:** Supervision, Writing – review & editing. **Tze Hui Liew:**
1249 Supervision, Project administration, Writing – review & editing. **Dip Nandi:**
1250 Supervision, Writing – review & editing.

1251 **Declaration of competing interest**

1252 The authors declare that they have no known competing financial interests
1253 or personal relationships that could have appeared to influence the work
1254 reported in this paper.

1255 **Funding**

1256 This research did not receive any specific grant from funding agencies in
1257 the public, commercial, or not-for-profit sectors.

1258 **Data availability**

1259 The data supporting the results of this study are openly available in the
1260 replication package (Katha et al., 2026), archived at Zenodo under <https://doi.org/10.5281/zenodo.23011153> and maintained in a public GitHub
1261 repository.³ It contains all measurements, the instrumentation, the analysis
1262 code and the scripts that generate every numeric value and every figure
1263 in this paper, including the C program and the script that builds each
1264 OpenSSL release, the exception baseline, and the cross-library harnesses in
1265 four languages with their pinned dependency manifests. Every number is
1266 generated from a file under the run or survey directories; none is entered by
1267 hand. The survey corpus is third-party source and is not redistributed; a
1268 manifest pinning every file examined by content hash stands in for it, covering
1269

³<https://github.com/Katha-Sikdar/A-Pre-Authentication-Key-Parse-Dominates-JWT-Validation-in-jsonwebtoken-and-Is-Load-Bearing->

1270 502 files across 457 repositories. Content hashes resolved for all of them;
1271 the repository HEAD commit resolved for 501 of 502, one repository having
1272 become unavailable between fetch and manifest. The content hash is the
1273 authoritative pin and is complete.

1274 **Ethics and responsible disclosure**

1275 This study measures software behaviour and publicly available source
1276 code; it involves no human participants and no personal data. No attack
1277 was run against any system the authors do not control. The key-confusion
1278 behaviour discussed in Section 8 is an existing, publicly documented attack
1279 class tested only through the library’s own test suite; the upstream report
1280 describes the constraint at the level of mechanism and does not include an
1281 exploitation recipe.

1282 **Declaration of generative AI and AI-assisted technologies in the** 1283 **manuscript preparation process**

1284 During the preparation of this manuscript, generative AI tools were used to
1285 improve language and readability. The authors reviewed all content following
1286 such use and take full responsibility for the contents of the manuscript.

1287 **Acknowledgements**

1288 The authors acknowledge ELITE Research Lab for providing the compu-
1289 tational resources and support necessary for this research.

1290 **References**

- 1291 Acar, Y., Backes, M., Fahl, S., Garfinkel, S., Kim, D., Mazurek, M.L., Stransky,
1292 C., 2017. Comparing the usability of cryptographic APIs, in: 2017 IEEE
1293 Symposium on Security and Privacy (SP). doi:10.1109/SP.2017.52.
- 1294 Auth0, 2022. CVE-2022-23539, CVE-2022-23541, CVE-2022-23540: Security
1295 update for jsonwebtoken. Auth0 security bulletin, 21 December 2022,
1296 [https://auth0.com/docs/secure/security-guidance/security-bul-](https://auth0.com/docs/secure/security-guidance/security-bulletins/2022-12-21-jwt-token)
1297 [letins/2022-12-21-jwt-token](https://auth0.com/docs/secure/security-guidance/security-bulletins/2022-12-21-jwt-token). Accessed 2026-09-28.

1298 Barrett, E., Bolz-Tereick, C.F., Killick, R., Mount, S., Tratt, L., 2017. Vir-
1299 tual machine warmup blows hot and cold. *Proceedings of the ACM on*
1300 *Programming Languages* 1, 52:1–52:27. doi:10.1145/3133876.

1301 Cabral, B., Marques, P., 2007. Exception handling: A field study in Java
1302 and .NET, in: *ECOOP 2007 – Object-Oriented Programming*, Springer.
1303 pp. 151–175. doi:10.1007/978-3-540-73589-2_8.

1304 Chen, J., Shang, W., 2017. An exploratory study of performance regression
1305 introducing code changes, in: *2017 IEEE International Conference on*
1306 *Software Maintenance and Evolution (ICSME)*. doi:10.1109/ICSME.2017
1307 .13.

1308 Cito, J., Schermann, G., Wittern, J.E., Leitner, P., Zumberi, S., Gall, H.C.,
1309 2017. An empirical analysis of the Docker container ecosystem on GitHub,
1310 in: *2017 IEEE/ACM 14th International Conference on Mining Software*
1311 *Repositories (MSR)*. doi:10.1109/MSR.2017.67.

1312 Coarfa, C., Druschel, P., Wallach, D.S., 2006. Performance analysis of TLS
1313 web servers. *ACM Transactions on Computer Systems* 24. doi:10.1145/11
1314 24153.1124155.

1315 Costa, D., Bezemer, C.P., Leitner, P., Andrzejak, A., 2021. What’s wrong
1316 with my benchmark results? Studying bad practices in JMH benchmarks.
1317 *IEEE Transactions on Software Engineering* 47, 1452–1467. doi:10.1109/
1318 TSE.2019.2925345.

1319 Curtsinger, C., Berger, E.D., 2013. STABILIZER: Statistically sound perfor-
1320 mance evaluation, in: *Proceedings of the 18th International Conference on*
1321 *Architectural Support for Programming Languages and Operating Systems*
1322 *(ASPLOS ’13)*, pp. 219–228. doi:10.1145/2451116.2451141.

1323 Davis, J.C., Williamson, E.R., Lee, D., 2018. A sense of time for JavaScript
1324 and Node.js: First-class timeouts as a cure for event handler poisoning, in:
1325 *Proceedings of the 27th USENIX Security Symposium*, pp. 343–359.

1326 Detering, D., Somorovsky, J., Mainka, C., Mladenov, V., Schwenk, J., 2017.
1327 On the (in-)security of JavaScript object signing and encryption, in: *Pro-*
1328 *ceedings of the 1st Reversing and Offensive-oriented Trends Symposium*
1329 *(ROOTS ’17)*. doi:10.1145/3150376.3150379.

1330 Egele, M., Brumley, D., Fratantonio, Y., Kruegel, C., 2013. An empirical
1331 study of cryptographic misuse in Android applications, in: Proceedings of
1332 the 2013 ACM SIGSAC Conference on Computer and Communications
1333 Security (CCS '13). doi:10.1145/2508859.2516693.

1334 Gaynor, A., Kehrer, P., 2026. The state of OpenSSL for pyca/cryptography.
1335 pyca/cryptography project statement, [https://cryptography.io/en/1](https://cryptography.io/en/latest/statements/state-of-openssl/)
1336 [atest/statements/state-of-openssl/](https://cryptography.io/en/latest/statements/state-of-openssl/). Maintainer statement, not peer
1337 reviewed; accessed 2026-09-24.

1338 Georges, A., Buytaert, D., Eeckhout, L., 2007. Statistically rigorous Java
1339 performance evaluation, in: Proceedings of the 22nd Annual ACM SIG-
1340 PLAN Conference on Object-Oriented Programming Systems, Languages
1341 and Applications (OOPSLA '07). doi:10.1145/1297027.1297033.

1342 Georgiev, M., Iyengar, S., Jana, S., Anubhai, R., Boneh, D., Shmatikov, V.,
1343 2012. The most dangerous code in the world: Validating SSL certificates
1344 in non-browser software, in: Proceedings of the 2012 ACM Conference on
1345 Computer and Communications Security (CCS '12). doi:10.1145/238219
1346 6.2382204.

1347 GitHub Advisory Database, 2015. Verification bypass in jsonwebtoken (CVE-
1348 2015-9235, GHSA-c7hr-j4mj-j2w6). [https://github.com/advisories/GH](https://github.com/advisories/GHSA-c7hr-j4mj-j2w6)
1349 [SA-c7hr-j4mj-j2w6](https://github.com/advisories/GHSA-c7hr-j4mj-j2w6). Accessed 2026-09-28.

1350 Gong, L., Pradel, M., Sen, K., 2015. JITProf: Pinpointing JIT-unfriendly
1351 JavaScript code, in: Proceedings of the 10th Joint Meeting on Foundations
1352 of Software Engineering (ESEC/FSE '15). doi:10.1145/2786805.2786831.

1353 Green, M., Smith, M., 2016. Developers are not the enemy! the need for usable
1354 security APIs. IEEE Security & Privacy 14. doi:10.1109/MSP.2016.111.

1355 Indela, S., Kulkarni, M., Nayak, K., Dumitras, T., 2016. Helping Johnny
1356 encrypt: Toward semantic interfaces for cryptographic frameworks, in:
1357 Proceedings of the 2016 ACM International Symposium on New Ideas, New
1358 Paradigms, and Reflections on Programming and Software (Onward! 2016).
1359 doi:10.1145/2986012.2986024.

1360 Jin, G., Song, L., Shi, X., Scherpelz, J., Lu, S., 2012. Understanding and
1361 detecting real-world performance bugs, in: Proceedings of the 33rd ACM

1362 SIGPLAN Conference on Programming Language Design and Implementa-
1363 tion (PLDI '12), pp. 77–88. doi:10.1145/2254064.2254075.

1364 Jones, M.B., Bradley, J., Sakimura, N., 2015. JSON Web Token (JWT).
1365 Request for Comments RFC 7519. Internet Engineering Task Force. doi:10
1366 .17487/RFC7519.

1367 Kalibera, T., Jones, R., 2013. Rigorous benchmarking in reasonable time, in:
1368 Proceedings of the 2013 International Symposium on Memory Management
1369 (ISMM '13). doi:10.1145/2464157.2464160.

1370 Kalliamvakou, E., Gousios, G., Blincoe, K., Singer, L., German, D.M., Damian,
1371 D., 2014. The promises and perils of mining GitHub, in: Proceedings of
1372 the 11th Working Conference on Mining Software Repositories (MSR 2014).
1373 doi:10.1145/2597073.2597074.

1374 Kaltenecker, C., Mühlbauer, S., Grebhahn, A., Siegmund, N., Apel, S., 2023.
1375 Performance evolution of configurable software systems: An empirical study.
1376 Empirical Software Engineering 28. doi:10.1007/s10664-023-10338-3.

1377 Katha, J.F., Prova, T.T., Morol, M.K., Liew, T.H., Nandi, D., 2026. Replica-
1378 tion package: a pre-authentication key parse dominates JWT validation in
1379 jsonwebtoken. doi:10.5281/zenodo.23011153. Source repository: [https://github.com/Katha-Sikdar/A-Pre-Authentication-Key-Parse-Dom-](https://github.com/Katha-Sikdar/A-Pre-Authentication-Key-Parse-Dominates-JWT-Validation-in-jsonwebtoken-and-Is-Load-Bearing-)
1380 [inates-JWT-Validation-in-jsonwebtoken-and-Is-Load-Bearing-](https://github.com/Katha-Sikdar/A-Pre-Authentication-Key-Parse-Dominates-JWT-Validation-in-jsonwebtoken-and-Is-Load-Bearing-).
1381 [inates-JWT-Validation-in-jsonwebtoken-and-Is-Load-Bearing-](https://github.com/Katha-Sikdar/A-Pre-Authentication-Key-Parse-Dominates-JWT-Validation-in-jsonwebtoken-and-Is-Load-Bearing-).

1382 Kery, M.B., Le Goues, C., Myers, B.A., 2016. Examining programmer
1383 practices for locally handling exceptions, in: Proceedings of the 13th
1384 International Conference on Mining Software Repositories (MSR '16), pp.
1385 484–487. doi:10.1145/2901739.2903497.

1386 Krüger, S., Späth, J., Ali, K., Bodden, E., Mezini, M., 2018. CrySL: An
1387 extensible approach to validating the correct usage of cryptographic APIs,
1388 in: 32nd European Conference on Object-Oriented Programming (ECOOP
1389 2018), pp. 10:1–10:27. doi:10.4230/LIPIcs.ECOOP.2018.10.

1390 Laaber, C., Leitner, P., 2018. An evaluation of open-source software micro-
1391 benchmark suites for continuous performance assessment, in: Proceedings
1392 of the 15th International Conference on Mining Software Repositories (MSR
1393 '18). doi:10.1145/3196398.3196407.

1394 Laaber, C., Scheuner, J., Leitner, P., 2019. Software microbenchmarking
1395 in the cloud. how bad is it really? Empirical Software Engineering 24.
1396 doi:10.1007/s10664-019-09681-1.

1397 Lazar, D., Chen, H., Wang, X., Zeldovich, N., 2014. Why does cryptographic
1398 software fail? A case study and open problems, in: Proceedings of the 5th
1399 Asia-Pacific Workshop on Systems (APSys '14). doi:10.1145/2637166.26
1400 37237.

1401 Leitner, P., Cito, J., 2016. Patterns in the chaos—a study of performance
1402 variation and predictability in public IaaS clouds. ACM Transactions on
1403 Internet Technology 16. doi:10.1145/2885497.

1404 Liu, Y., Xu, C., Cheung, S.C., 2014. Characterizing and detecting performance
1405 bugs for smartphone applications, in: Proceedings of the 36th International
1406 Conference on Software Engineering (ICSE '14), pp. 1013–1024. doi:10.1
1407 145/2568225.2568229.

1408 Malka, J., Zacchiroli, S., Zimmermann, T., 2026. Docker does not guarantee
1409 reproducibility. arXiv preprint arXiv:2601.12811.

1410 Mytkowicz, T., Diwan, A., Hauswirth, M., Sweeney, P.F., 2009. Producing
1411 wrong data without doing anything obviously wrong!, in: Proceedings of the
1412 14th International Conference on Architectural Support for Programming
1413 Languages and Operating Systems (ASPLOS '09). doi:10.1145/1508244.
1414 1508275.

1415 Mytkowicz, T., Diwan, A., Hauswirth, M., Sweeney, P.F., 2010. Evaluating
1416 the accuracy of Java profilers, in: Proceedings of the 31st ACM SIGPLAN
1417 Conference on Programming Language Design and Implementation (PLDI
1418 '10), pp. 187–197. doi:10.1145/1806596.1806618.

1419 Nadi, S., Krüger, S., Mezini, M., Bodden, E., 2016. Jumping through hoops, in:
1420 Proceedings of the 38th International Conference on Software Engineering
1421 (ICSE '16). doi:10.1145/2884781.2884790.

1422 Naser, A., Sadeghi Lahijani, M., Wu, C., Gavahi, M., Hoang, V.T., Wang,
1423 Z., Yuan, X., 2020. Performance evaluation and modeling of cryptographic
1424 libraries for MPI communications. arXiv preprint arXiv:2010.06139.

1425 Nistor, A., Jiang, T., Tan, L., 2013a. Discovering, reporting, and fixing
1426 performance bugs, in: Proceedings of the 10th Working Conference on
1427 Mining Software Repositories (MSR '13). doi:10.1109/MSR.2013.6624035.

1428 Nistor, A., Song, L., Marinov, D., Lu, S., 2013b. Toddler: Detecting perfor-
1429 mance problems via similar memory-access patterns, in: Proceedings of
1430 the 35th International Conference on Software Engineering (ICSE '13), pp.
1431 562–571. doi:10.1109/ICSE.2013.6606602.

1432 Ogasawara, T., Komatsu, H., Nakatani, T., 2006. EDO: Exception-directed
1433 optimization in Java. ACM Transactions on Programming Languages and
1434 Systems 28. doi:10.1145/1111596.1111598.

1435 de Pádua, G.B., Shang, W., 2017. Studying the prevalence of exception
1436 handling anti-patterns, in: Proceedings of the 25th International Conference
1437 on Program Comprehension (ICPC '17), pp. 328–331. doi:10.1109/ICPC
1438 .2017.1.

1439 Parravicini, A., Mueller, T., 2021. The cost of speculation: Revisiting over-
1440 heads in the V8 JavaScript engine, in: 2021 IEEE International Symposium
1441 on Workload Characterization (IISWC). doi:10.1109/IISWC53511.2021.
1442 00013.

1443 Rahaman, S., Xiao, Y., Afrose, S., Shaon, F., Tian, K., Frantz, M., Kantar-
1444 cioglu, M., Yao, D.D., 2019. CryptoGuard: High precision detection of
1445 cryptographic vulnerabilities in massive-sized Java projects, in: Proceedings
1446 of the 2019 ACM SIGSAC Conference on Computer and Communications
1447 Security (CCS '19), pp. 2455–2472. doi:10.1145/3319535.3345659.

1448 Renwick, J., Spink, T., Franke, B., 2019. Low-cost deterministic C++ ex-
1449 ceptions for embedded systems, in: Proceedings of the 28th International
1450 Conference on Compiler Construction (CC 2019). doi:10.1145/3302516.
1451 3307346.

1452 Selakovic, M., Glaser, T., Pradel, M., 2017. An actionable performance
1453 profiler for optimizing the order of evaluations, in: Proceedings of the
1454 26th ACM SIGSOFT International Symposium on Software Testing and
1455 Analysis (ISSTA '17). doi:10.1145/3092703.3092716.

1456 Selakovic, M., Pradel, M., 2016. Performance issues and optimizations in
1457 JavaScript: An empirical study, in: Proceedings of the 38th International

1458 Conference on Software Engineering (ICSE '16), pp. 61–72. doi:10.1145/
1459 2884781.2884829.

1460 Shatnawi, A., Al-Duwairi, B., Smarneh, A., 2024. Comprehensive empirical
1461 study of Python JWT libraries. *Procedia Computer Science* 238, 827–832.
1462 doi:10.1016/j.procs.2024.06.099.

1463 Sheffer, Y., Hardt, D., Jones, M.B., 2020. JSON Web Token Best Current
1464 Practices. Request for Comments RFC 8725. Internet Engineering Task
1465 Force. doi:10.17487/RFC8725.

1466 Sikeridis, D., Kampanakis, P., Devetsikiotis, M., 2020. Assessing the overhead
1467 of post-quantum cryptography in TLS 1.3 and SSH, in: *Proceedings of the*
1468 *16th International Conference on Emerging Networking EXperiments and*
1469 *Technologies (CoNEXT '20)*. doi:10.1145/3386367.3431305.

1470 Song, L., Lu, S., 2014. Statistical debugging for real-world performance
1471 problems, in: *Proceedings of the 2014 ACM International Conference*
1472 *on Object Oriented Programming Systems Languages & Applications*
1473 *(OOPSLA '14)*, pp. 561–578. doi:10.1145/2660193.2660234.

1474 Staicu, C.A., Pradel, M., 2018. Freezing the web: A study of ReDoS vul-
1475 nerabilities in JavaScript-based web servers, in: *Proceedings of the 27th*
1476 *USENIX Security Symposium*, pp. 361–376.

1477 Stefan, P., Horký, V., Bulej, L., Tůma, P., 2017. Unit testing performance in
1478 Java projects: Are we there yet?, in: *Proceedings of the 8th ACM/SPEC*
1479 *International Conference on Performance Engineering (ICPE '17)*, pp. 401–
1480 412. doi:10.1145/3030207.3030226.

1481 Traini, L., Cortellessa, V., Di Pompeo, D., Tucci, M., 2023. Towards effective
1482 assessment of steady state performance in Java software: Are we there yet?
1483 *Empirical Software Engineering* 28. doi:10.1007/s10664-022-10247-x.

1484 Tuveri, N., Brumley, B.B., 2019. Start your ENGINes: Dynamically loadable
1485 contemporary crypto, in: *2019 IEEE Cybersecurity Development (SecDev)*.
1486 doi:10.1109/SecDev.2019.00014.

1487 Weimer, W., Necula, G.C., 2004. Finding and preventing run-time error
1488 handling mistakes, in: *Proceedings of the 19th Annual ACM SIGPLAN*

1489 Conference on Object-Oriented Programming, Systems, Languages, and
1490 Applications (OOPSLA '04). doi:10.1145/1028976.1029011.

1491 Xu, B., Jia, S., Lin, J., Zheng, F., Ma, Y., Liu, L., Gu, X., Song, L.,
1492 2023. JWTKey: Automatic cryptographic vulnerability detection in JWT
1493 applications, in: Computer Security – ESORICS 2023, Springer. doi:10.1
1494 007/978-3-031-51479-1_14.

1495 Yang, J., Wang, E., Chen, J., Wang, Q., Zhang, Y., Duan, H., Xie, W.,
1496 Wang, B., 2026. Token time bomb: Evaluating JWT implementations for
1497 vulnerability discovery, in: Proceedings of the Network and Distributed
1498 System Security Symposium (NDSS '26).

1499 Zaman, S., Adams, B., Hassan, A.E., 2012. A qualitative study on performance
1500 bugs, in: Proceedings of the 9th IEEE Working Conference on Mining
1501 Software Repositories (MSR '12), pp. 199–208. doi:10.1109/MSR.2012.6
1502 224281.

1503 Zerouali, A., Mens, T., Robles, G., Gonzalez-Barahona, J.M., 2019. On
1504 the relation between outdated Docker containers, severity vulnerabilities,
1505 and bugs, in: Proceedings of the 26th IEEE International Conference on
1506 Software Analysis, Evolution and Reengineering (SANER '19), pp. 491–501.
1507 doi:10.1109/SANER.2019.8668013.

1508 Zhu, X., She, G., Xue, B., Zhang, Y., Zhang, Y., Zou, X.K., Duan, X., He,
1509 P., Krishnamurthy, A., Lentz, M., Zhuo, D., Mahajan, R., 2023. Dissecting
1510 overheads of service mesh sidecars, in: Proceedings of the 2023 ACM
1511 Symposium on Cloud Computing (SoCC '23). doi:10.1145/3620678.3624
1512 652.